Blocking and Filtering Techniques for Entity Resolution: A Survey

GEORGE PAPADAKIS, University of Athens, Greece
DIMITRIOS SKOUTAS, IMSI, Athena Research Center, Greece
EMMANOUIL THANOS, KU Leuven, Belgium
THEMIS PALPANAS, Paris Descartes University, France

Entity Resolution (ER), a core task of Data Integration, detects different entity profiles that correspond to the same real-world object. Due to its inherently quadratic complexity, a series of techniques accelerate it so that it scales to voluminous data. In this survey, we review a large number of relevant works under two different but related frameworks: Blocking and Filtering. The former restricts comparisons to entity pairs that are more likely to match, while the latter identifies quickly entity pairs that are likely to satisfy predetermined similarity thresholds. We also elaborate on hybrid approaches that combine different characteristics. For each framework we provide a comprehensive list of the relevant works, discussing them in the greater context. We conclude with the most promising directions for future work in the field.

1 INTRODUCTION

Entity Resolution (ER) is the task of identifying different entity profiles that describe the same real-world object [29, 47]. It is a core task for Data Integration, applying to any kind of data, from the structured entities of relational databases [24] to the semistructured entities of the Linked Open Data Cloud (https://lod-cloud.net) [29, 38] and the unstructured entities that are automatically extracted from free text [148]. ER consists of two parts: (1) the *candidate selection step*, which determines the entities worth comparing, and (2) the *candidate matching step*, or simply *Matching*, which compares the selected entities to determine whether they represent the same real-world

object. The latter step involves *pairwise comparisons*, i.e., time-consuming operations that typically apply string similarity measures to pairs of entities, dominating the overall cost of ER [24, 29, 38].

In this survey, we focus on the candidate selection step, which is the crucial part of ER with respect to time efficiency and scalability. Without it, ER suffers from a quadratic time complexity, $O(n^2)$, as every entity profile has to be compared with all others. Reducing this computational cost is the goal of numerous techniques from two dominant frameworks: Blocking and Filtering. The former attempts to identify entity pairs that are likely to match, restricting comparisons only between them, while the latter attempts to quickly discard pairs that are guaranteed to not match, executing comparisons only between the rest. The former operates without knowledge of the Matching step, while the latter is based on it, assuming that two entities match if their similarity exceeds a specified threshold. Hence, Blocking and Filtering share the same goal but are complementary, as they operate under different settings and assumptions. So far, though, they have been developed independently of one another: their combination and, more generally, their relation have been overlooked in the literature, with the exception of very few works (e.g., [82]).

Moreover, the rise of Big Data poses new challenges for both Blocking and Filtering approaches [29, 38]: *Volume* requires techniques to scale to millions of entities, while *Variety* calls for techniques that can cope with an unprecedented schema heterogeneity. Both Blocking and Filtering address Volume primarily through parallelization. Existing techniques were adapted to split their workload into smaller chunks that are distributed across different processing units so that they are executed in parallel. This can be done on a cluster (distributed methods) or through the modern multicore and multisocket hardware architectures. Variety, though, is addressed differently in each field. For Blocking, the schema-aware methods are replaced by schema-agnostic techniques, which disregard any schema information, creating blocks of very high recall but low precision. Additionally, a whole new category of methods, called *Block Processing*, intervenes between Blocking and Matching to refine the original blocks, significantly increasing precision at a negligible (if any) cost in recall. For Filtering, techniques that employ more relaxed matching criteria (e.g., fuzzy set matching or local string similarity join) are proposed, while the case of low-similarity thresholds is also considered.

To the best of our knowledge, this is the first survey to comprehensively cover the aforementioned aspects and to jointly review the two frameworks for efficient ER. We formally define Blocking, Block Processing, and Filtering, introducing a common terminology that facilitates their understanding. For each field, we propose a new taxonomy with categories that highlight the distinguishing characteristics of the corresponding methods. Based on these taxonomies, we provide a broad overview of every field, elucidating the functionality of the main techniques as well as the relations among them. As a result, established techniques are now seen in a different light; Canopy Clustering [97], for instance, may now be viewed as a Block Processing method. We also elaborate on the parallelization methods for each field. Most importantly, this survey attempts to place Blocking and Filtering under a common context, taking special care to stress hybrid methods that combine features from both Blocking and Filtering, to analyze works that experimentally compare the two frameworks (e.g., [154]) and to qualitatively outline their commonalities and differences. We also investigate the ER tools that incorporate established efficiency techniques and propose a series of open challenges that constitute promising directions for future research.

Parts of the material included in this survey have been presented in tutorials at WWW 2014 [157], ICDE 2016 [123], ICDE 2017 [156], and WWW 2018 [124]. A past survey [25] also covers efficiency ER techniques but is restricted to the schema-aware Blocking methods. Other surveys [47] and textbooks [24, 29, 38] provide a holistic overview of ER, merely examining the main Blocking and Block Processing techniques. Closer to our work is a recent survey on Blocking [111], which, however, offers a more limited coverage and refers neither to parallelization nor to Filtering

works. Recent surveys on string and set similarity joins also exist, but they focus exclusively on centralized [64, 95, 187] or distributed approaches [50], with the purpose of experimental comparison, and without covering approximate techniques that allow for more relaxed matching criteria. Most importantly, none of these surveys considers similarity joins in the broader context of ER.

The rest of the article is structured as follows: Section 2 provides background knowledge on ER and its efficiency techniques, while Sections 3 and 4 delve into Blocking and Block Processing, respectively. Section 5 is devoted to Filtering, whereas Section 6 elaborates on works that combine Blocking with Filtering. Section 8 enumerates the main ER tools that incorporate efficiency methods, Section 7 provides a high-level discussion of the relation between Blocking and Filtering, Section 9 provides the main directions for future work, and Section 10 concludes the paper.

2 PRELIMINARIES

At the core of ER lies the notion of *entity profile*, which constitutes a uniquely identified description of a real-world object in the form of name-value pairs. Assuming infinite sets of attribute names $\mathcal{N}$, attribute values $\mathcal{V}$, and unique identifiers $\mathcal{I}$, an entity profile is formally defined as follows [29, 120]:

Definition 1 (Entity Profile). An *entity profile* $\mathbf{e_{id}}$ is a tuple $\langle id, A_{id} \rangle$, where $id \in \mathcal{I}$ is a unique identifier, and A_{id} is a set of name-value pairs $\langle n, v \rangle$, with $n \in \mathcal{N}$ and $v \in (\mathcal{V} \cup \mathcal{I})$. A set of entity profiles $\mathcal{E}$ is called *entity collection*.

This definition is simple but flexible enough to accommodate a wide variety of (semi-)structured representations; e.g., nested attributes can be transformed into a flat set of name-value pairs, while links may be represented by assigning the id of one entity as the attribute value of the other.

Definition 2 (Entity Resolution). Two entity profiles e_i and e_j *match*, $\mathbf{e_i} \equiv \mathbf{e_j}$, if they refer to the same real-world entity. Matching entities are also called *duplicates*. The task of Entity Resolution (ER) is to find all matching entities within an entity collection or across two or more entity collections.

In particular, we distinguish between the following two cases [24, 25]:

(1) *Deduplication* receives as input an entity collection $\mathcal{E}$ and produces as output the set of all pairs of matching entity profiles within $\mathcal{E}$, i.e., $\mathcal{D}(\mathcal{E}) = \{(e_i, e_j) : e_i \in \mathcal{E}, e_j \in \mathcal{E}, e_i \equiv e_j\}$.

(2) *Record Linkage* receives two duplicate-free entity collections, $\mathcal{E}_1$ and $\mathcal{E}_2$, and returns the pairs of matching entity profiles between them, i.e., $\mathcal{D}(\mathcal{E}_1, \mathcal{E}_2) = \{(e_i, e_j) : e_i \in \mathcal{E}_1, e_j \in \mathcal{E}_2, e_i \equiv e_j\}$.

Multisource Entity Resolution involves three or more entity collections and can be performed by applying Deduplication to the union of all collections or by executing a sequence of pairwise Record Linkage tasks, provided that every input collection is duplicate-free.

ER performance is characterized by its *effectiveness* and its *efficiency*. The former refers to how many of the actual duplicates are detected, while the latter expresses the computational cost for detecting them—usually in terms of the number of performed comparisons, which is referred to as *cardinality* and denoted by $||\mathcal{E}||$. The naive, brute-force approach performs all pairwise comparisons between the input entity profiles, having a quadratic complexity that does not scale to large datasets; for Record Linkage, $||\mathcal{E}|| = |\mathcal{E}_1| \times |\mathcal{E}_2|$, while for Deduplication, $||\mathcal{E}|| = |\mathcal{E}| \cdot (|\mathcal{E}| - 1)/2$.

Blocking. To tackle ER's inherently quadratic complexity, Blocking trades slightly lower effectiveness for significantly higher efficiency. Its goal is to reduce the number of performed comparisons, while missing as few matches as possible. Ideally, one would compare only the pairs of duplicates, whose number grows *linearly,* with the number of the input entity profiles [53, 156].

To this end, Blocking clusters potentially matching entities in common blocks and exclusively compares entity profiles that co-occur in at least one block.

Internally, a blocking method employs a *blocking scheme*, which applies to one or more entity collections to yield a set of blocks $\mathcal{B}$, called *block collection*. Cardinality $||\mathcal{B}||$ denotes the number of comparisons in $\mathcal{B}$, given that only entity pairs within the same block are compared, i.e., $||\mathcal{B}|| = \sum_{b_i \in \mathcal{B}} ||b_i||$, where $||b_i||$ stands for the number of comparisons contained in an individual block b_i. We denote the set of *detectable duplicates* in $\mathcal{B}$ as $\mathcal{D}(\mathcal{B})$, while $\mathcal{D}(\mathcal{E})$ stands for all existing duplicates. Since $\mathcal{B}$ reduces the number of performed comparisons, $\mathcal{D}(\mathcal{B}) \subseteq \mathcal{D}(\mathcal{E})$.

A common assumption in the literature is the *oracle*, i.e., a perfect matching function that, for each pair of entity profiles, decides correctly whether they match or not [25, 38, 120, 121, 156]. Using an oracle, a pair of duplicates is detected as long as they share at least one block. This allows for reasoning about the performance of blocking methods independently of matching methods: there is a clear tradeoff between the effectiveness and the efficiency of a blocking scheme [25, 38, 156]: the more comparisons are contained in the resulting block collection $\mathcal{B}$ (i.e., higher $||\mathcal{B}||$), the more duplicates will be detected (i.e., higher $|\mathcal{D}(\mathcal{B})|$), raising effectiveness at the cost of lower efficiency. Thus, a blocking scheme should achieve a good balance between these two competing objectives as expressed through the following measures [16, 33, 100, 116]:

(1) *Pair Completeness (PC)* corresponds to *recall*, estimating the portion of the detectable duplicates in $\mathcal{B}$ with respect to those in $\mathcal{E}$: $PC(\mathcal{B}) = |\mathcal{D}(\mathcal{B})|/|\mathcal{D}(\mathcal{E})| \in [0, 1]$.

(2) *Pair Quality (PQ)* corresponds to *precision*, estimating the portion of comparisons in $\mathcal{B}$ that correspond to real duplicates: $PQ(\mathcal{B}) = |\mathcal{D}(\mathcal{B})|/||\mathcal{B}|| \in [0, 1]$.

(3) *Reduction Ratio (RR)* measures the reduction in the number of pairwise comparisons in $\mathcal{B}$ with respect to the brute-force approach: $RR(\mathcal{B}, \mathcal{E}) = 1 - ||\mathcal{B}||/||\mathcal{E}|| \in [0, 1]$.

Higher values for *PC* indicate higher *effectiveness* of the blocking scheme, while higher values for *PQ* and *RR* indicate higher *efficiency*. Note that *PC* provides an optimistic estimation of recall, presuming the existence of an oracle, while *PQ* provides a pessimistic estimation of precision, treating as false positives the repeated comparisons between duplicates (i.e., only the nonrepeated duplicate pairs are considered as true positives). In this context, we can define Blocking as follows:

Definition 3 (Blocking). Given an entity collection $\mathcal{E}$, Blocking clusters similar entities into a block collection $\mathcal{B}$ such that $PC(\mathcal{B})$, $PQ(\mathcal{B})$, and $RR(\mathcal{B}, \mathcal{E})$ are simultaneously maximized.

This definition refers to Deduplication but can be easily extended to Record Linkage. Simultaneously maximizing *PC*, *PQ*, and *RR* necessitates that the enhancements in efficiency do not affect the effectiveness of Blocking, carefully removing comparisons between nonmatching entities. Conceptually, Blocking can be viewed as an optimization task, but this implies that the real duplicate collection $\mathcal{D}(\mathcal{E})$ is known, which is actually what ER tries to compute. Hence, Blocking is typically treated as an engineering task that provides an approximate solution for the data at hand.

A blocking-based ER workflow may comprise several stages. First, *Block Building* (BlBu) applies a blocking scheme to produce a block collection $\mathcal{B}$ from the input entity collection(s). This step may be repeated several times on the same input, applying multiple blocking schemes, in order to achieve a more robust performance in the context of highly noisy data. Often, there is a second, optional stage, called *Block Processing*, which refines $\mathcal{B}$ through additional optimizations that further reduce the number of performed comparisons. This may involve discarding *entire blocks* that primarily contain unnecessary comparisons, called *Block Cleaning* (BlCl), and/or discarding *individual comparisons* within certain blocks, called *Comparison Cleaning* (CoCl). The former may be applied repeatedly, each time enforcing a different, complementary method to discard blocks,
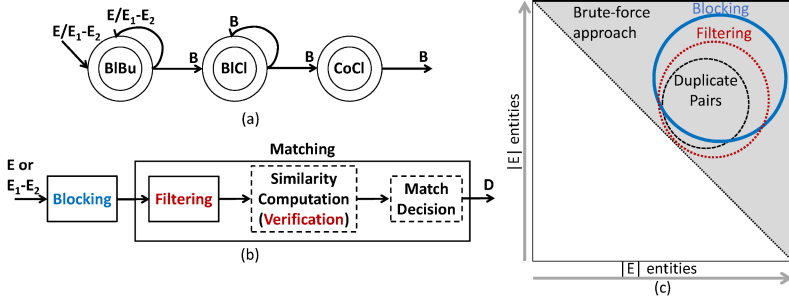
Fig. 1. (a) The internal functionality of Blocking modeled as a deterministic finite automaton with three states: Block Building (BlBu), Block Cleaning (BlCl), and Comparison Cleaning (CoCl). (b) The end-to-end workflow for nonlearning Entity Resolution [82]. (c) The relative computational cost for the brute-force approach, Blocking, Filtering, and the ideal solution (Duplicate Pairs) over Deduplication.

but the latter can be performed only once; CoCl comprises competitive methods that serve exactly the same purpose and, once applied to a block collection, they alter it in such a way that turns all other methods inapplicable. Figure 1(a) models this workflow as a deterministic finite automaton with three states, where each state corresponds to one of the blocking subtasks.

Filtering. Given two entity collections $\mathcal{E}_1$ and $\mathcal{E}_2$, a similarity function $f_S : \mathcal{E}_1 \times \mathcal{E}_2 \to \mathbb{R}$, and a similarity threshold θ, a *similarity join* identifies all pairs of entity profiles in $\mathcal{E}_1$ and $\mathcal{E}_2$ that have similarity at least θ, i.e., $\mathcal{E}_1 \bowtie_\theta \mathcal{E}_2 = \{(e_i, e_j) \in \mathcal{E}_1 \times \mathcal{E}_2 : f_S(e_i, e_j) \geq \theta\}$.

Similarity joins can be used for defining ER under the intuitive assumption that matching entity profiles are highly similar. In fact, the above formulation corresponds to Record Linkage, while Deduplication can be defined analogously as a self-join operation, where $\mathcal{E}_1 \equiv \mathcal{E}_2$.

To avoid exhaustive pairwise comparisons, similarity joins typically follow the *filter-verification* framework, which involves two stages [10, 64]:

(1) *Filtering* computes a set of *candidates* for each entity e_i, excluding all those that cannot match with e_i. In other words, it prunes all true negatives but allows some false positives.
(2) *Verification* computes the actual similarity between candidates (or a sufficient upper bound) to remove the false positives.

Due to the relatively straightforward implementation of Verification, in the following we exclusively focus on Filtering. The relevant techniques are defined with respect to three parameters: (i) the representation for each entity, (ii) the similarity function between entity pairs under this representation, and (iii) the similarity threshold above which two entities are considered to match.

The representation typically relies on *signatures* extracted from each entity such that two entities match only if their signatures overlap. Given that we address ER over entities described by one or more textual attributes, we focus on string similarity joins, which can be *character* or *token based*. The former compare two strings by representing them as sequences of characters and by considering the character transformations required to transform one string into the other. The latter are also called *set similarity joins*, since they transform the strings into sets, typically via tokenization or q-gram extraction, and then compare strings using a set-based similarity measure.

Regarding the similarity function, the most common one for character-based similarity joins is Edit Distance, which measures the minimum number of edit operations (i.e., insertions, deletions, and substitutions) that are required to transform one string to the other [10]. For token-based similarity joins, the most commonly used similarity measures include Overlap, Jaccard, Cosine, or Dice. The last three are normalized variants of the Overlap [10, 64, 95].

Type of Joins	Measure	Definition	Equivalent Overlap Threshold
character-based	Edit Distance	# character transformations	$\max(|x|, |y|)+1-(1+\theta) \times q$
token-based	Overlap	$|x \cap y|$	θ
	Cosine	$|x \cap y|/\sqrt{|x| \cdot |y|}$	$\theta \times \sqrt{|x| \cdot |y|}$
	Dice	$2 \cdot |x \cap y|/(|x| + |y|)$	$\theta \times (|x| + |y|)/2$
	Jaccard	$|x \cap y|/(|x| + |y| - |x \cap y|)$	$\theta \times (|x| + |y|)/(1 + \theta)$

Fig. 2. Definition of the main similarity measures used by string similarity join algorithms, and how the input threshold θ for each measure can be transformed into an equivalent Overlap threshold τ.

Finally, the similarity threshold depends on the data at hand. Note, though, that the join algorithms do not operate directly with thresholds on Jaccard, Cosine, or Dice similarity, but first translate the given threshold θ into an equivalent set overlap threshold τ that depends on the size of the sets, as shown in Figure 2. A similar transformation is also possible for Edit Distance, which means that set similarity joins may be applied to this measure as well [10].

Blocking vs. Filtering. The relation between the two frameworks is illustrated in Figure 1(b). Blocking, in the sense of the entire process in Figure 1(a), is applied first, reducing the pairwise comparisons that are considered by Matching. These comparisons are further cut down by Filtering, which is subsequently applied, as the initial part of Matching, given that it requires specifying both a similarity measure and a similarity threshold. Next, Verification is applied to estimate the actual similarity between the compared attribute values. The Entity Resolution process concludes with *Match Decision*, which synthesizes the estimated similarity between multiple attribute values to determine whether the compared entity profiles are indeed duplicates.

Both Blocking and Filtering are optional steps, but at least one of them should be applied in order to tame the otherwise quadratic computational cost of ER. As shown in Figure 1(c), Blocking yields a *superlinear* but *subquadratic* time complexity, lying between the two extremes: the brute-force solution and the ideal one (i.e., Duplicate Pairs). The same applies to the computational cost of Filtering, except that it typically constitutes an *exact* procedure that produces no false negatives, i.e., missed duplicates. It exclusively allows false positives, which are later removed by Verification [10]. For this reason, Filtering corresponds to a superset of Duplicate Pairs in Figure 1(c). In contrast, Blocking constitutes an inherently *approximate* solution that increases ER efficiency at the cost of allowing both false positives and false negatives [29]. Thus, it intersects Duplicate Pairs, such that the area of their intersection is inversely proportional to the duplicates that are missed by Blocking, while the relative complement of the Duplicate Pairs in Blocking is analogous to the executed comparisons between nonmatching entities.

Note that Figure 1(c) corresponds to Deduplication but can be easily generalized to Record Linkage as well. Moreover, the relative performance of Blocking and Filtering, i.e., the relative position of their circles, depends on the methods and the data at hand. In most cases, though, the best solution is to use both frameworks, yielding the computational cost that corresponds to their intersection. However, this approach is rarely used in the literature (e.g., [82]). Most works on Blocking typically omit Filtering (e.g., [25, 112, 128]), whereas most works on Filtering disregard Blocking, applying directly to the input entity collections (e.g., [64, 95]). The goal of the present survey is to cover this gap, elucidating the complementarity of the two frameworks.

3   BLOCK BUILDING

Block Building receives as input one or more entity collections and produces as output a block collection $\mathcal{B}$. The process is guided by a *blocking scheme*, which determines how entity profiles are assigned to blocks. This scheme typically comprises two parts. First, every entity is processed

to extract *signatures* (e.g., tokens), such that the similarity of signatures reflects the similarity of the corresponding profiles. Second, every entity is mapped to one or more blocks based on these signatures. Let $\mathcal{P}(\mathcal{S})$ denote the power set of a set S and $\mathcal{K}$ denote the universe of signatures appearing in entity profiles. We formally define a blocking scheme as follows:

Definition 4 (Blocking Scheme). Given an entity collection $\mathcal{E}$, a *blocking scheme* is a function $f_B : \mathcal{E} \to \mathcal{P}(\mathcal{B})$ that maps entity profiles to blocks. It is composed of two functions: (1) a *transformation* function $f_T : \mathcal{E} \to \mathcal{P}(\mathcal{K})$ that maps an entity profile to a set of *signatures* (also called *blocking keys*) and (2) an *assignment* function $f_A : \mathcal{K} \to \mathcal{P}(\mathcal{B})$ that maps each signature to one or more blocks.

This definition applies to Deduplication but can be easily extended to Record Linkage.

The set of comparisons in the resulting block collection $\mathcal{B}$ is called *comparison collection* and is denoted by $C(\mathcal{B})$. Every comparison $c_{i,j} \in C(\mathcal{B})$ belongs to one of the following types [120, 121]:

- *Matching comparison*, if e_i and e_j match
- *Superfluous comparison*, if e_i and e_j do not match
- *Redundant comparison*, if e_i and e_j have already been compared in a previous block

We collectively call the last two types *unnecessary comparisons*, as their execution brings no gain.

Note that the resulting block collection $\mathcal{B}$ can be modeled as an inverted index that points from block ids to entity ids. For this reason, Block Building is also called *Indexing* [24, 25].

3.1 Taxonomy

To facilitate the understanding of the main methods for Block Building, we organize them into a novel taxonomy that consists of the following dimensions:

- *Key selection* distinguishes between *nonlearning* and *learning-based* methods. The former rely on rules derived from expert knowledge or mere heuristics, while the latter require a training set to learn the best blocking keys using Machine Learning techniques.
- *Schema awareness* distinguishes between *schema-aware* and *schema-agnostic* methods. The former extract blocking keys from specific attributes that are considered to be more appropriate for matching (e.g., more distinctive or less noisy), while the latter disregard schema knowledge, extracting blocking keys from all attributes.
- *Key type* distinguishes between *hash-* or *equality-based* methods, which map a pair of entities to the same block if they have a common key, and *sort-* or *similarity-based* methods, which map a pair of entities to the same block if they have a similar key. There exist also *hybrid* methods, which combine hash- with sort-based functionality.
- *Redundancy awareness* classifies methods into three categories based on the relation between their blocks. *Redundancy-free* methods assign every entity to a single block, thus creating disjoint blocks. *Redundancy-positive* methods place every entity into multiple blocks, yielding overlapping blocks. The more blocks two entities share, the more similar their profiles are. The number of blocks shared by a pair of entities is thus proportional to their matching likelihood. *Redundancy-neutral* methods create overlapping blocks, where most entity pairs share the same number of blocks, or the degree of redundancy is arbitrary, having no implications.
- *Constraint awareness* distinguishes blocking methods into *lazy*, which impose no constraints on the blocks they create, and *proactive*, which enforce constraints on their blocks (e.g., maximum block size), or refine their comparisons by discarding unnecessary ones.

Table 1. Taxonomy of the Block Building Methods Discussed in Sections 3.2 and 3.3.

Method	Key Type	Redundancy Awareness	Constraint Awareness	Matching Awareness
Standard Blocking (SB) [49]	hash based	redundancy free	lazy	static
Suffix Arrays Blocking (SA) [3]	hash based	redundancy positive	proactive	static
Extended Suffix Arrays Blocking [25, 112]	hash based	redundancy positive	proactive	static
Improved Suffix Arrays Blocking [33]	hash based	redundancy positive	proactive	static
Q-Grams Blocking [25, 112]	hash based	redundancy positive	lazy	static
Extended Q-Grams Blocking [11, 25, 112]	hash based	redundancy positive	lazy	static
MFIBlocks [75]	hash based	redundancy positive	proactive	static
Sorted Neighborhood (SN) [60, 61, 132]	sort based	redundancy neutral	proactive	static
Extended Sorted Neighborhood [25]	sort based	redundancy neutral	lazy	static
Incrementally Adaptive SN [185]	sort based	redundancy neutral	proactive	static
Accumulative Adaptive SN [185]	sort based	redundancy neutral	proactive	static
Duplicate Count Strategy (DCS) [41]	sort based	redundancy neutral	proactive	dynamic
DCS++ [41]	sort based	redundancy neutral	proactive	dynamic
Sorted Blocks [40]	hybrid	redundancy neutral	lazy	static
Sorted Blocks New Partition [40]	hybrid	redundancy neutral	proactive	static
Sorted Blocks Sliding Window [40]	hybrid	redundancy neutral	proactive	static
(a) Nonlearning, schema-aware methods				
ApproxRBSetCover [16]	hash based	redundancy positive	lazy	static
ApproxDNF [16]	hash based	redundancy positive	lazy	static
Blocking Scheme Learner (BSL) [100]	hash based	redundancy positive	lazy	static
Conjunction Learner [21] (semi-supervised)	hash based	redundancy positive	lazy	static
BGP [48]	hash based	redundancy positive	lazy	static
CBlock [146]	hash based	redundancy positive	proactive	static
DNF Learner [54]	hash based	redundancy positive	lazy	dynamic
FisherDisjunctive [72] (unsupervised)	hash based	redundancy positive	lazy	static
(b) Learning-based (supervised), schema-aware methods				
Token Blocking (TB) [116]	hash based	redundancy positive	lazy	static
Attribute Clustering Blocking [120]	hash based	redundancy positive	lazy	static
RDFKeyLearner [154]	hash based	redundancy positive	lazy	static
Prefix-Infix(-Suffix) Blocking [119]	hash based	redundancy positive	lazy	static
TYPiMatch [92]	hash based	redundancy positive	lazy	static
Semantic Graph Blocking [109]	-	redundancy neutral	proactive	static
(c) Nonlearning, schema-agnostic methods				
Hetero [73]	hash based	redundancy positive	lazy	static
Extended DNF BSL [74]	hash based	redundancy positive	lazy	static
(d) Learning-based (unsupervised), schema-agnostic methods				

- *Matching awareness* distinguishes between *static* methods, which are independent of the subsequent matching process, producing an immutable block collection, and *dynamic* methods, which intertwine Block Building with Matching, updating or processing their blocks dynamically, as more duplicates are detected.

Table 1 maps all methods discussed in Sections 3.2 and 3.3 to our taxonomy.
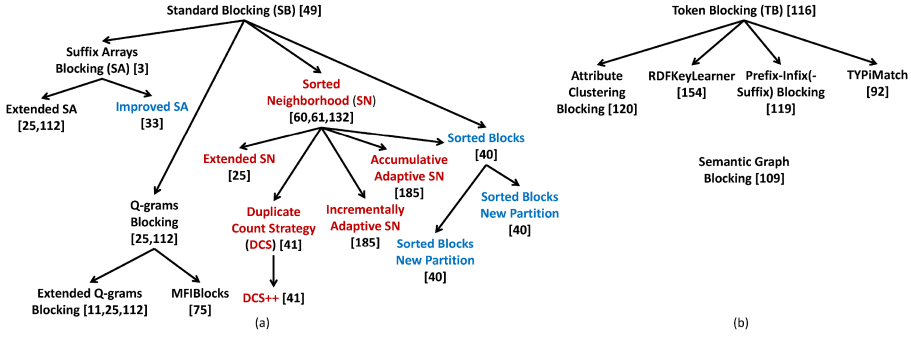
Fig. 3. The genealogy trees of nonlearning (a) schema-aware and (b) schema-agnostic Block Building techniques. Hybrid, hash-, and sort-based methods are marked in blue, black, and red, respectively.

3.2 Schema-Aware Block Building

Methods of this type assume that the input entity profiles adhere to a known schema, and based on this schema and respective domain knowledge, one can select the attributes that are most suitable for Blocking. We distinguish between nonlearning methods, reviewed in Section 3.2.1, and learning-based methods, reviewed in Section 3.2.2.

3.2.1 Nonlearning Methods. The family tree of the methods in this category is shown in Figure 3(a); a parent-child edge implies that the latter method improves upon the former one. Below, we elaborate on these methods based on their key type.

Hash-Based Methods. *Standard Blocking* (SB) [49] involves the simplest functionality: an expert selects the most suitable attributes, and a transformation function concatenates (parts of) their values to form blocking keys. For every distinct key, a block is created containing all corresponding entities. In short, SB operates as a hash function, conveying two main advantages: (1) it yields redundancy-free blocks, and (2) it has a linear time complexity, $O(|E|)$. On the flip side, its effectiveness is very sensitive to noise, as the slightest difference in the blocking keys of duplicates places them in different blocks. SB is also a lazy method that imposes no limit on block sizes.

To address these issues, *Suffix Arrays Blocking* (SA) [3] converts each blocking key of SB into the list of its suffixes that are longer than a predetermined minimum length l_{min}. Then, it defines a block for every suffix that does not exceed a predetermined frequency threshold b_{max}, which essentially specifies the maximum block size. This proactive functionality is necessary, as very frequent suffixes (e.g., "ing") result in large blocks that are dominated by unnecessary comparisons.

SA has two major advantages [33]: (1) it has low time complexity, $O(|E| \cdot log|E|)$ [4], and is very efficient, as it results in a small but relevant set of candidate matches, and (2) it is very effective, due to the robustness to the noise at the beginning of blocking keys and the high levels of redundancy (i.e., it places every entity into multiple blocks). On the downside, SA does not handle noise at the end of SB keys; e.g., two matches with SB keys "JohnSnith" and "JohnSmith" have no common suffix if $l_{min} = 4$, while for $l_{min} = 3$, they co-occur in a block only if the frequency of "ith" is lower than b_{max}.

This problem is addressed by *Extended Suffix Arrays Blocking* [25, 112], which uses as keys not just the suffixes of SB keys, but all their substrings with more than l_{min} characters; e.g., for $l_{min} = 4$, SA extracts from "JohnSnith" the keys "JohnSnith," "ohnSnith," "hnSnith," "nSnith," "Snith," and "nith," while *Extended SA* additionally extracts the keys "John," "ohnS," "hnSn," "nSni," and "Snit," as well as all substrings of "JohnSnith" ranging from five to eight characters.

Another extension of SB is *Q-grams Blocking* [25, 112]. Its transformation function converts the blocking keys of SB into subsequences of q characters (*q-grams*) and defines a block for every distinct q-gram. For example, for $q = 3$, the key *france* is transformed into the trigrams *fra*, *ran*, *anc*, *nce*. This approach differs from Extended SA in that it does not restrict block sizes (lazy method). Also, it is more resilient to noise than SB but results in more and larger blocks.

To improve it, *Extended Q-Grams Blocking* [11, 25, 112] uses combinations of q-grams instead of individual q-grams. Its transformation function concatenates at least l q-grams, where $l = max(1, \lfloor k \cdot t \rfloor)$, with k denoting the number of q-grams and $t \in [0, 1)$ standing for a user-defined threshold. The larger t is, the larger l gets, yielding fewer keys from the k q-grams. For $T = 0.9$ and $q = 3$, the key *france* is transformed into the following four signatures ($k = 4$ and $l = 3$): [*fra*, *ran*, *anc*, *nce*], [*fra*, *ran*, *anc*], [*fra*, *anc*, *nce*], [*ran*, *anc*, *nce*]. In this way, q-gram-based blocking keys become more distinctive, decreasing the number and cardinality of blocks.

A more advanced q-gram-based approach is *MFIBlocks* [75]. Its transformation function concatenates keys of Q-Grams Blocking into itemsets and uses a maximal frequent itemset algorithm to define as new blocking keys those exceeding a predetermined support threshold.

Sort-Based Methods. *Sorted Neighborhood* (SN) [60] sorts all blocking keys in alphabetical order and arranges the associated entities accordingly. Subsequently, a window of fixed size w slides over the sorted list of entities and compares the entity at the last position with all other entities placed within the same window. The underlying assumption is that the closer the blocking keys of two entities are in the lexicographical order, the more likely they are to be matching. Originally crafted for relational data, SN is extended to hierarchical/XML data based on user-defined keys in [132].

SN has three major advantages [25]: (1) it has low time complexity, $O(|E| \cdot log|E|)$; (2) it results in linear ER complexity, $O(w \cdot |E|)$; and (3) it is robust to noise, supporting errors at the end of blocking keys. However, it may place two entities in the same block even if their keys are dissimilar (e.g., "alphabet" and "apple," if no other key intervenes between them). Its performance also depends heavily on the window size w, which is difficult to configure, especially in Deduplication, where the matching entities form clusters of varying size [25, 40].

To ameliorate the effect of w, a common solution is the *Multi-pass SN* [61], which applies the core algorithm multiple times, using a different transformation function in each iteration. In this way, more matches can be identified, even if the window is set to low size. Another solution is the *Extended Sorted Neighborhood* [25, 112], which slides a window of fixed size over the sorted list of blocking keys rather than the list of entities; this means that each block merges w SB blocks.

More advanced strategies adapt the window size dynamically to optimize the balance between effectiveness and efficiency. They are grouped into three categories, depending on the criterion for moving the boundaries of the window [91]:

1) *Key similarity strategy.* The window size increases if the similarity of the blocking keys exceeds a predetermined threshold, which indicates that more similar entities should be expected [91].

2) *Entity similarity strategy.* The window size relies on the similarity of the entities within the current window. *Incrementally Adaptive SN* [185] increases the window size if the distance of the first and the last element in the window is smaller than a predetermined threshold. The actual increase depends on the current window size and the selected threshold. *Accumulative Adaptive SN* [185] creates windows with a single overlapping entity and exploits transitivity to group multiple adjacent windows into the same block, as long as the last entity of one window is a potential duplicate of the last entity in the next window. After

expanding the window, both algorithms apply a retrenchment phase that decreases the window size until all entities are potential duplicates.

3) *Dynamic strategy.* The core assumption is that the more duplicates are found within a window, the more are expected to be found by increasing its size. *Duplicate Count Strategy* (DCS) [41] defines a window w for every entity in SN's sorted list and executes all its comparisons to compute the ratio d/c, where d denotes the newly detected duplicates and c the executed comparisons. The window size is then incremented by one position at a time as long as $d/c \geq \phi$, where $\phi \in (0, 1)$ is a threshold that expresses the average number of duplicates per comparison. DCS++ [41] improves DCS by increasing the window size with the next $w - 1$ entities, even if the new ratio becomes lower than ϕ. Using transitive closure, it skips some windows, saving part of the comparisons.

Hybrid Methods. *Sorted Blocks* [40] combines the benefits of SB and SN. First, it sorts all blocking keys and the corresponding entities in lexicographical order, like SN. Then, it partitions the sorted entities into disjoint blocks, like SB, using a prefix of the blocking keys. Next, all pairwise comparisons are executed within each block. To avoid missing any matches, an overlap parameter o defines a window of fixed size that includes the o last entities in the current block together with the first entity of the next block. The window slides by one position at a time until reaching the o^{th} entity of the next block, executing all pairwise comparisons between entities from different blocks.

Sorted Blocks is a lazy approach that does not restrict block sizes. Thus, it may result in large blocks that dominate its processing time. To address this, two proactive variants set a limit on the maximum block size. *Sorted Blocks New Partition* [40] creates a new block when the maximum size is reached for a (prefix of) blocking key; the overlap between the blocks ensures that every entity is compared with its predecessors and successors in the sorting order. *Sorted Blocks Sliding Window* [40] avoids executing all comparisons within a block that is larger than the upper limit; instead, it slides a window equal to the maximum block size over the entities of the current block.

Finally, *Improved Suffix Arrays Blocking* [33] employs the same blocking keys as SA but sorts them in alphabetical order, like SN. Then, it compares the consecutive keys with a string similarity measure. If the similarity of two suffixes exceeds a predetermined threshold, the corresponding blocks are merged in an effort to detect duplicates even when there is noise at the end of SB keys or their sole common key is too frequent. For example, *Improved SA* detects the high string similarity of the keys "JohnSnith" and "JohnSmith," placing the corresponding entities into the same block.

3.2.2 Learning-Based Methods. We distinguish these methods into supervised and unsupervised ones. Both rely on a labeled dataset that includes pairs of matching and nonmatching entities, called *positive* and *negative instances*, respectively. This dataset is used to learn *blocking predicates*, i.e., combinations of an attribute name and a transformation function (e.g., $\{title, First3Characters\}$). Entities sharing the same output for a particular blocking predicate are considered candidate matches (i.e., hash-based functionality). Disjunctions of conjunctions of predicates, i.e., composite blocking schemes, are learned by optimizing an objective function.

Supervised Methods. *ApproxRBSetCover* [16] learns disjunctive blocking schemes by solving a standard weighted set cover problem. The cover is iteratively constructed by adding in each turn the blocking predicate that maximizes the ratio of the previously uncovered positive pairs over the covered negative pairs. This is a "soft cover," since some positive instances may remain uncovered.

ApproxDNF [16] alters ApproxRBSetCover so that it learns blocking schemes in Disjunctive Normal Form (DNF). Instead of individual predicates, each turn greedily learns a conjunction of up to k predicates that maximizes the ratio of positive and negative covered instances.

A similar approach is *Blocking Scheme Learner* (BSL) [100]. Based on an adaptation of the Sequential Covering Algorithm, it learns blocking schemes that maximize *RR*, while maintaining *PC* above a predetermined threshold. Its output is a disjunction of conjunctions of blocking predicates.

BSL is improved by *Conjunction Learner* [21], which minimizes the candidate matches not only in the labeled but also in the *unlabeled* data, while maintaining high *PC*. The effect of the unlabeled data is determined through a weight $w \in [0, 1]$; $w = 0$ disregards unlabeled data completely, falling back to BSL, while $w = 1$ indicates that they are as equally important as the labeled ones.

In another line of research, *Blocking based on Genetic Programming* (BGP) [48] employs a tree representation of supervised blocking schemes, where every leaf node corresponds to a blocking predicate. In every turn, a set of genetic programming operators, such as copy, mutation, and crossover, are applied to the initial, random set of blocking schemes. Then, a fitness function infers the performance of the new schemes from the harmonic mean of *PC* and *RR*, and the best ones are returned as output. Yet, BGP involves numerous internal parameters that are hard to fine-tune.

Another tree-based approach is *CBLOCK* [146]. In this case, every edge is annotated with a hash (i.e., transformation) function and every node n_i comprises the set of entities that result after applying all hash functions from the root to n_i. *CBLOCK* is the only proactive learning-based method, restricting the maximum size of its blocks. Every node that exceeds this limit is split into smaller, disjoint blocks through a greedy algorithm that picks the best hash function based on the resulting *PC*. To minimize the human effort, a drill down approach is proposed for bootstrapping.

Unsupervised Methods. *FisherDisjunctive* [72] uses a weak training set generated by the TF-IDF similarity between pairs of entities. Pairs with very low (high) values are considered as negative (positive) instances. A Boolean feature vector is then associated with every labeled instance. The discovery of DNF blocking schemes is finally cast as a Fisher feature selection problem.

Similarly, *DNF Learner* [54] applies a matching algorithm to a sample of entity pairs to automatically create a labeled dataset. Then, the learning of blocking schemes is cast as a DNF learning problem. To scale it to the exponential search space of possible schemes, their complexity is restricted to manageable levels (e.g., they comprise at most k predicates).

3.3 Schema-Agnostic Block Building

Methods of this type make no assumptions about schema knowledge, disregarding completely attribute names; they extract blocks from all attribute values. Thus, they inherently support noise in both attribute names and values and are suitable for highly heterogeneous, loosely structured entity profiles, such as those stemming from the Web of Data [116, 119, 120].

Nonlearning Methods. The family tree of this category appears in Figure 3(b). The cornerstone approach is *Token Blocking* (TB) [116]. Assuming that duplicates share at least one common token, its transformation function extracts all tokens from all attribute values of every entity. A block b_t is then defined for every distinct token t. Hence, two entities co-occur in block b_t if they share token t in their values, regardless of the associated attribute names.

To improve TB, *Attribute Clustering Blocking* [120] requires the common tokens of two entities to appear in *syntactically similar attributes*. These are attribute names that correspond to similar values but are not necessarily semantically matching (unlike Schema Matching). First, it clusters attributes based on the similarities of their aggregate values. Each attribute is connected to its most similar one and the transitive closure of the connected attributes forms disjoint clusters. A block $b_{k,t}$ is then defined for every token t in the values of the attributes belonging to cluster k.

RDFKeyLearner [154] applies TB independently to the values of specific attributes, which are selected through the following process: each attribute is associated with a *discriminability* score, which amounts to the portion of distinct values over all values in the given dataset. If this is lower than a predetermined threshold, the attribute is ignored due to limited diversity; i.e., too

many entities have the same value(s). For each attribute with high discriminability, its *coverage* is estimated, i.e., the portion of entities that contain it. The harmonic mean of discriminability and coverage is then computed for all valid attributes and the one with the maximum score is selected for defining blocking keys as long as its score exceeds another predetermined threshold. If not, the selected attribute is combined with all other attributes and the process is repeated.

Prefix-Infix(-Suffix) Blocking [119] exploits the naming pattern in entity URIs. The *prefix* describes the domain of the URI, the *infix* is a local identifier, and the optional *suffix* contains details about the format, or a named anchor [114]; e.g., in the URI `https://en.wikipedia.org/wiki/France#History`, the prefix is `https://en.wikipedia.org/wiki`, the infix is `France`, and the suffix is `History`. In this context, this method uses as keys all (URI) infixes along with all tokens in the literal values.

TYPiMatch [92] improves TB by automatically detecting the entity types in the input data. It creates a co-occurrence graph, where every node corresponds to a token in any attribute value and every edge connects two tokens if both conditional probabilities of co-occurrence exceed a predetermined threshold. The maximal cliques are extracted and merged if their overlap exceeds another threshold. The resulting clusters correspond to the entity types, with every entity participating in all types to which its tokens belong. TB is then applied independently to the profiles of each type.

Finally, *Semantic Graph Blocking* [109] is based exclusively on the relations between entities, be it foreign keys in a database or links in RDF data. It completely disregards attribute values, building a collaborative graph, where every node corresponds to an entity and every edge connects two associated entities. For instance, the collaborative graph for a bibliographic data collection can be formed by mapping every author to a node and adding edges between coauthors. Then a new block b_i is formed for each node n_i, containing all nodes connected with n_i through a path, provided that the path length or the block size does not exceed predetermined limits (proactive functionality).

Learning-Based Methods. *Hetero* [73] converts the input data into heterogeneous structured datasets using property tables. Then, it maps every entity to a normalized TF vector and applies an adapted Hungarian algorithm with linear scalability to produce positive and negative feature vectors. Finally, it applies *FisherDisjunctive* [72] with bagging to achieve robust performance.

Similarly, *Extended DNF BSL* [74] combines an established instance-based schema matcher with weighted set covering to learn DNF blocking schemes with at most k predicates.

3.4 Parallelization Approaches

To scale Blocking methods to massive entity collections without altering their functionality, the *MapReduce framework* [35] is typically used, as it offers fault-tolerant, optimized execution for applications distributed across a set of independent nodes.

Schema-Aware Methods. The hash-based, nonlearning methods are adapted to MapReduce in a straightforward way. The `map` phase implements the transformation function(s), emitting (`key`, `entity_id`) pairs for each entity. Every reducer acts as an assignment function, placing all entities with blocking key t in block b_t. Dedoop [77] provides such implementations for various methods.

For sort-based methods, the adaptation of SN to MapReduce in [79] can be used as a template. The `map` function extracts the blocking key(s) from each input entity, while the ensuing *partitioning* phase sorts all entities in alphabetical order of their keys. The `reduce` function slides a window of fixed size within every reduce partition. Inevitably, entities close to the partition boundaries need to be compared across different reduce tasks. Thus, the `map` function is extended to replicate those entities, forwarding them to the respective reduce task and its successor.

DCS and DCS++ are adapted to the MapReduce framework in [99], using three jobs. The first one sorts the originally unordered entities of the data partition assigned to each mapper according

to the selected blocking keys. It also selects the boundary pairs of the sorted partitions. The second job generates the Partition Allocation Matrix, which specifies the sorted partitions to be replicated, while the third job performs DCS(++) locally, to the data assigned to every reducer.

Schema-Agnostic Methods. A single MapReduce job is required for parallelizing TB [29, 46]. For every input entity e_i, the map function emits a (t, e_i) pair for every token t in the values of e_i. Then, all entities sharing a particular token are directed to the same reducer to form a new block.

For Attribute Clustering Blocking, four MapReduce jobs are required [29, 46]. The first assembles all values per attribute. The second computes the pairwise similarities between all attributes, even if they are placed in different data partitions. The third connects every attribute to its most similar one. The fourth associates every attribute name with a cluster id and adapts TB's map function to emit pairs of the form $(k.t, i)$, where k is the cluster id of e_i's attribute name that contains token t.

Finally, the parallelization of Prefix-Infix(-Suffix) Blocking involves three MapReduce jobs [29, 46]. The first parallelizes the algorithm that extracts the prefixes from a set of URIs [114]. The second extracts the URI suffixes. The third applies TB's mapper to the literal values simultaneously with an infix mapper that emits a pair (j, e_i) for every infix j that is extracted from e_i's profile. The final reduce phase ensures that all entities having a common token or infix are placed in the same block.

Load Balancing. For MapReduce, it is crucial to distribute evenly the overall workload among the available nodes, avoiding potential bottlenecks. The following methods distribute the execution of comparisons in a block collection—not the cost of building the blocks.

BlockSplit [78] partitions large blocks into smaller subblocks and processes them in parallel. Every entity is compared to all entities in its subblock as well as to all entities of its superblock, even if their subblock is initially assigned to a different node. This yields an additional network and I/O overhead and may still lead to unbalanced workload, due to subblocks of different size.

To overcome this, *PairRange* [78] splits evenly the comparisons in a set of blocks into a predefined number of partitions. It involves a single MapReduce job with a mapper that associates every entity e_i in block b_k with the output key $p.k.i$, where p denotes the partition id. The reducer assembles all entities that have the same p and block id, reproducing the comparisons of each partition.

The space requirements of these two algorithms are improved in [186], which minimizes their memory consumption by adapting them so that they work with sketches.

Finally, *Dis-Dedup* [30] is the only method that takes into account both the computational and the communication cost (e.g., network transfer time, local disk I/O time). Dis-Dedup considers all possible cases, from disjoint blocks produced by a single blocking technique to overlapping blocks derived from multiple techniques. It also provides strong theoretical guarantees that the overall maximum cost per reducer is within a small constant factor from the lower bounds.

3.5 Discussion and Experimental Results

The performance of the above techniques is examined both qualitatively and quantitatively in a series of individual of papers (e.g., [34, 110, 116, 120]) and experimental analyses (e.g., [25, 112, 128]). Below, we summarize the main findings in order to facilitate the use of Block Building techniques.

Starting with *Standard Blocking* (SB), its performance depends heavily on the frequency distribution of attribute values and, thus, of blocking keys. The best case corresponds to a uniform distribution, where $||B|| = ||E||/|B|$ [25]. Due to its lazy functionality, though, all other key distributions yield a portion of large blocks with many superfluous comparisons, i.e., low *PQ* and *RR*.

Suffix Arrays Blocking (SA) improves SB's *PC* by supporting errors at the beginning of blocking keys. The higher l_{min} is and the lower b_{max} is, the lower $||B||$ and *PC* get. For the same settings, *Extended SA* raises *PC* at the cost of higher $||B||$, which inevitably lowers both *PQ* and *RR*. *Improved*

SA is theoretically proven in [34] to result in a *PC* greater than or equal to that of SA, though at the cost of a higher computational cost and more comparisons, which lower *PQ* and *RR*.

Q-grams Blocking yields higher *PC* than SB but decreases both *PQ* and *RR*. *Extended Q-grams Blocking* raises *PQ* and *RR* at a limited, if any, cost in *PC*. MFIBlocks reduces significantly the number of blocks and matching candidates (i.e., very high *PQ* and *RR*) [75], but it may come at the cost of missed matches (insufficient *PC*) in case the resulting blocking keys are very restrictive for matches with noisy descriptions [128].

For *Sorted Neighborhood* (SN), a small *w* leads to high *PQ* and *RR* but low *PC* and vice versa for a large *w*. For *Extended SN*, variations in the window size have a large impact on efficiency (*PQ* and *RR*), affecting the portion of unnecessary comparisons, but *PC* is more stable. Among the other SN variants, DCS++ is theoretically proven to miss no matches with an appropriate value for ϕ, while being at least as efficient as SN. *Sorted Blocks New Partition* outperforms most SN-based algorithms but includes more parameters than SN, involving a more complex configuration.

Most importantly, all these nonlearning schema-aware methods are quite parameter sensitive: even small-parameter-value modifications may yield significantly different performance [25, 33, 110, 112]. Their most important parameter is the definition of the blocking keys, which requires fine-tuning by an expert. Otherwise, their *PC* remains insufficient, placing most duplicates in no common block [25, 112]. This applies even to methods that employ redundancy for higher recall.

This shortcoming is ameliorated by schema-agnostic methods, which consistently achieve much higher *PC* than their schema-aware counterparts [112]. They also simplify the configuration of Block Building, reducing its sensitivity through the automatic definition of blocking keys [112, 128]. Rather than human intervention or expert knowledge, their robustness emanates from the high levels of redundancy they employ, placing every entity in a multitude of blocks. On the down-side, they yield a considerably higher number of comparisons, resulting in very low *PQ* and *RR*. Both, however, can be significantly improved by Block Processing [122, 128].

Regarding the relative performance of schema-agnostic methods, TB yields very high *PC*, at the cost of very low *PQ* and *RR*. It constitutes a very efficient approach, iterating only once over the input entities, and it is the sole parameter-free Block Building technique in the literature as well as the most generic one, applying to any entity collection with textual values. Its performance is improved by *Attribute Clustering* and *Prefix-Infix(-Suffix) Blocking* for specific types of datasets: highly heterogeneous ones, with a large variety of attribute names [120, 128], and semistructured (RDF) ones [116], respectively. In these cases, both methods yield a much larger number of smaller blocks, significantly raising *PQ* at a minor cost in *PC*. Both methods, though, involve a much higher computational cost than TB. The same applies to *TYPiMatch*, where the detection of entity types is a rather time-consuming process. Yet, its *PC* is consistently insufficient, because it falsely divides duplicate entities into different entity types, due to the sensitivity to its parameter configuration [128].

Finally, the learning-based Block Building techniques typically suffer from the scarcity of labelled datasets; even if a training set is available for a particular dataset, it cannot be directly used for learning supervised blocking schemes for another dataset. Instead, a complex transfer learning procedure is typically required [103, 164]. Regarding their efficiency, BSL is typically faster than *ApproxRBSetCover* and *ApproxDNF*, as it exclusively considers positive instances, thus requiring a smaller training set. *Conjunction Learner* requires every supervised blocking scheme to be applied to the large set of unlabeled data, which is impractical. To accelerate it, a random sample of the unlabeled data is used in practice. *CBLOCK* is also the only learning-based method that is suitable for the MapReduce framework: every entity runs through the learned tree and is directed to the machine corresponding to its leaf node. In terms of effectiveness, there is no clear winner. BSL and

FisherDisjunctive achieve the top performance in [110]. The latter addresses the scarcity of labeled data but is not scalable to large datasets.

4 BLOCK PROCESSING

Block Processing receives as input an existing block collection $\mathcal{B}$ and produces as output a new block collection $\mathcal{B}'$ that improves the balance between effectiveness and efficiency, i.e., $PQ(B) \ll PQ(B')$, $RR(B', B) \gg 0$, while $PC(B) \sim PC(B')$. We distinguish Block Processing methods into *Block Cleaning* ones, which decide whether entire blocks should be retained or modified, and *Comparison Cleaning* ones, which decide whether individual comparisons are unnecessary.

4.1 Block Cleaning

We classify Block Cleaning methods into two categories: (1) *static*, which are independent of matching results, and (2) *dynamic*, which are interwoven with the matching process.

Static Methods. A core idea is the assumption that the larger a block is, the less likely it is to contain unique duplicates, i.e., matches that share no other block. Such large blocks are typically produced by lazy schema-agnostic techniques and correspond to stop words. In this context, *Block Purging* discards blocks that exceed an upper limit on block cardinality [120] or size [119]. *Block Filtering* [127] applies this assumption to individual entities, removing every entity from the largest blocks that contain it. In other words, it retains every entity in $r\%$ of its smallest blocks.

In a different line of research, *Size-based Block Clustering* [51] applies hierarchical clustering to transform a set of blocks into a new one where all block sizes lie within a specified size range. It merges recursively small blocks that correspond to similar blocking keys, while splitting large blocks into smaller ones. A penalty function controls the tradeoff between block quality and block size. A similar approach is the MapReduce-based dynamic blocking algorithm in [98], which splits large blocks into subblocks. *MaxIntersectionMerge* [102] ensures that all blocks involve at least $|b|_{min}$ entities. To this end, it merges each block smaller than $|b|_{min}$ entities with the block that has the most entities in common and is larger than $|b|_{min}$. Similarly, *Rollup Canopies* [146] receives as input a training set with positive examples, a limit on the maximum block size, and a set of disjoint blocks; using a learning-based greedy algorithm, it merges pairs of small blocks to increase *PC*.

Finally, [139] generalizes Meta-blocking (see Section 4.2) to Multi-source ER: it constructs a graph, where the nodes correspond to blocks and the edges connect blocks whose blocking keys are more similar than a predetermined threshold. The edges are weighted using various functions and all pairs of blocks are then processed in descending edge weights in an effort to maximize the redundant and superfluous comparisons that are skipped.

Dynamic Methods. *Iterative Blocking* [179] merges any new pair of detected duplicates, e_i and e_j, into a new entity, $e_{i,j}$, and replaces both e_i and e_j with $e_{i,j}$ in all blocks that contain them, even if they have already been processed. The new entity $e_{i,j}$ is compared with all co-occurring entities, as the new content in $e_{i,j}$ might identify previously missed matches. The ER process terminates when all blocks have been processed without finding new duplicates.

Iterative Blocking applies exclusively to Deduplication. In Record Linkage, there is no need for merging two matching entities, due to the 1-1 restriction. Still, the detected duplicates should be propagated in order to save the superfluous comparisons with their co-occurring entities in the subsequently processed blocks. The earlier the matches are detected, the more superfluous comparisons are saved. To this end, *Block Scheduling* optimizes the processing order of blocks in a noniterative way, sorting them in decreasing order of the probability $p_i(d)$ that a block b_i contains a pair of duplicates. This is set inversely proportional to block cardinality, i.e., $p_i(d) = 1/||b_i||$ [152], or to the minimum size of the inner block, i.e., $p_i(d) = 1/min|b_{i,1}|, |b_{i,2}|$, where $|b_{i,k}| \subset \mathcal{E}_k$ [116].
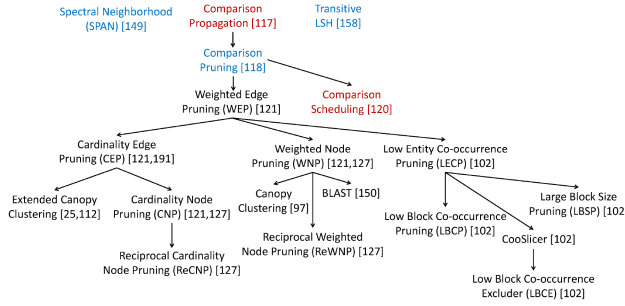
Fig. 4. The genealogy tree of nonlearning Comparison Cleaning methods. Methods in black conform to the Meta-blocking framework in Figure 5, methods in blue are Meta-blocking techniques following a (partially) different approach, and methods in red are not part of the Meta-blocking framework.


Fig. 5. (a) A block collection B with $e_1 \equiv e_3$ and $e_2 \equiv e_4$, (b) the corresponding Entity Index, (c) the corresponding blocking graph G_B, (d) the weighted G_B, (e) the pruned G_B, and (f) the new block collection B'.

The former definition also applies to Iterative Blocking, which does not specify the exact block processing order, even though this affects significantly the resulting performance [128].

Block Pruning [116] extends Block Scheduling by exploiting the decreasing density of detected matches in its processing order (i.e., the later a block is processed, the less unique duplicates it contains). After processing the latest block, it estimates the average number of executed comparisons per new duplicate. If this ratio falls below a specific threshold, it terminates the ER process.

4.2 Comparison Cleaning

Nonlearning Methods. Figure 4 illustrates the family tree of the methods belonging to this category. The cornerstone method is *Comparison Propagation* [117], which propagates all executed comparisons to the subsequently processed blocks. In this manner, it eliminates all redundant comparisons in a given block collection without losing any pair of duplicates, thus raising *PQ* and *RR* at no cost in *PC*. It builds an inverted index that points from entity ids to block ids, called *Entity Index*, and with its help, it compares two entities e_i and e_j in block b_k only if k is their least common block id. For example, consider the blocks in Figure 5(a) and their Entity Index in Figure 5(b). The least common block id of e_1 and e_3 is 2. Thus, they are compared in b_2, but neither in b_4 nor in b_5.

Given a redundancy-positive block collection, the Entity Index allows for identifying the blocks shared by a pair of co-occurring entities. This allows for weighting all pairwise comparisons in proportion to the matching likelihood of the corresponding entities, based on the principle that the more blocks two entities share, the more likely they are to be matching. This gives rises to a family of *Meta-blocking* techniques [121, 127, 150] that go beyond Comparison Propagation by discarding not only all redundant comparisons, but also the vast majority of the superfluous ones.

The first relevant method is *Comparison Pruning* [118], which computes the Jaccard coefficient of the block lists of two entities. If it does not exceed a conservative threshold that depends on the average number of blocks per entity, the comparison is pruned, as it designates an unlikely match.

Meta-blocking was formalized into a more principled approach in [121]. The given redundancy-positive block collection $\mathcal{B}$ is converted into a blocking graph G_B, where the nodes correspond to entities and the edges connect every pair of co-occurring entities (see Figure 5(c)). Given that no parallel edges are allowed, all redundant comparisons are discarded by definition. The edges are then weighted proportionally to the likelihood that the adjacent entities are matching. In Figure 5(d), the edge weights indicate the number of common blocks. Edges with low weights are pruned, because they correspond to superfluous comparisons. In Figure 5(e), all edges with a weight lower than the average one are discarded. The resulting pruned blocking graph $G_{B'}$ is transformed into a restructured block collection $\mathcal{B}'$ by forming one block for every retained edge (see Figure 5(f)). As a result, $\mathcal{B}'$ exhibits a much higher efficiency, $PQ(B') \gg PQ(B)$ and $RR(B', B) \gg 0$, for similar effectiveness, $PC(B') \sim PC(B)$; in our example, the 12 comparisons in the input blocks of Figure 5(a) are reduced to two matching comparisons in the output blocks in Figure 5(f).

Four main pruning algorithms exist: (1) *Weighted Edge Pruning* (WEP) removes all edges that do not exceed a specific threshold, e.g., the average edge weight [121]; (2) *Cardinality Edge Pruning* (CEP) retains the globally top-K top weighted edges, where K is static [121] or dynamic [191]; (3) *Weighted Node Pruning* (WNP) retains in each node neighborhood the entities that exceed a local threshold, which may be the average edge weight of each neighborhood [121], or the average of the maximum weights in the two adjacent node neighborhoods, as in *BLAST* [150]; (4) *Cardinality Node Pruning* (CNP) retains the k weighted edges in each node neighborhood [121]. *Reciprocal WNP* and *CNP* [127] apply an aggressive pruning that retains edges satisfying the pruning criteria in both adjacent node neighborhoods. WNP and WEP are combined through the weighted sum of their thresholds in [9].

Another family of pruning algorithms is presented in [102], focusing on the edge weights between the entities in each block. *Low Entity Co-occurrence Pruning* (LECP) cleans every block from a specific portion of the entities with the lowest average edge weights. *Large Block Size Pruning* (LBSP) applies LECP only to the blocks whose size exceeds the average block size in the input block collection. *Low Block Co-occurrence Pruning* (LBCP) removes every entity from the blocks, where it is connected with the lowest weights, on average, with the rest of the entities. *CooSlicer* enforces a maximum block size constraint, $|b|_{max}$, to all input blocks. In blocks larger than $|b|_{max}$, all entities are sorted in decreasing order of average edge weight, and the $|b|_{max}$ top-ranked entities are iteratively placed into a new block. *Low Block Co-occurrence Excluder* (LBCE) discards a specific portion of the blocks with the lowest average edge weight among their entities.

All these pruning algorithms can be coupled with any *edge-weighting scheme* [121]. ARCS sums the inverse cardinalities of the common blocks, giving higher weights to entity pairs that co-occur in smaller blocks. CBS counts the number of blocks shared by two entities, as in Figure 5(c), with ECBS extending it to discount the contribution from entities placed in many blocks. JS corresponds to the Jaccard coefficient of two block lists, while EJS extends it to discount the contribution from entities appearing in many nonredundant comparisons. Finally, Pearson's χ^2 test assesses whether two adjacent entities appear independently in blocks and can be combined with the aggregate attribute entropy associated with the tokens forming their common blocks [150].

Note that Meta-blocking covers established methods that are considered as Block Building methods in the literature: given that Block Building is equivalent to indexing [25], any method based on indexes is in fact a Meta-blocking technique. For example, *Transitive LSH* [158] converts the blocks extracted from LSH into an unweighted blocking graph and applies a community detection algorithm (e.g., [31]) to partition the graph nodes into disjoint clusters, which will become the new blocks. The process finishes when the size of the largest cluster is lower than a predetermined threshold. This approach can be applied on top of any Block Building method, not just LSH.

The generalization principle also applies to *Canopy Clustering* [97], which places all entities in a pool, and in every iteration, it removes a random entity e_i from the pool to create a new block. Using a cheap similarity measure, all entities still in the pool are compared with e_i. Those exceeding a threshold t_{ex} are removed from the pool and placed into the new block. Entities exceeding another threshold t_{in} ($<t_{ex}$) are also placed in the new block, without being removed from the pool. As the cheap similarity measure, we can use any of the above weighting schemes on top of any Block Building method, thus turning Canopy Clustering into a pruning algorithm for Meta-blocking.

The generalization applies to *Extended Canopy Clustering* [25, 112], too, which replaces the sensitive weight thresholds with cardinality ones: for each randomly selected entity, the n_1 nearest entities are placed in its block, while the n_2($\leq n_1$) nearest entities are removed from the pool.

In another line of research, *SPAN* [149] converts a block collection into a matrix M, where the rows correspond to entities and the columns to the tf-idf of blocking keys (tokens or q-grams). Then, the entity-entity matrix is defined as $A = MM^T$. A spectral clustering algorithm converts A into a binary tree, where the root node contains all entities and every leaf node is a disjoint subset of entities. The Newman-Girvan modularity is used as the stopping criterion for the bipartition of the tree. Blocks are then derived from a search procedure that carries out pairwise comparisons based on the blocking keys, inside the leaf nodes, and across the neighboring ones.

Finally, the sole dynamic nonlearning method is *Comparison Scheduling* [120]. Its goal is to detect most matches upfront so as to maximize the superfluous comparisons that are skipped, due to the 1-1 restriction. It orders all comparisons in decreasing matching likelihood (edge weight) and executes a comparison only if none of the involved entities has already been matched.

Learning-Based Approaches. *Supervised Meta-blocking* [125] treats edge pruning as a binary classification problem, where every edge is labeled "`likely match`" or "`unlikely match`." Every edge is represented by a feature vector that comprises five features: ARCS, ECBS, JS, and the Node Degrees of the adjacent entities. Undersampling is employed to tackle the class imbalance problem: the training set makes up just 5% of the minority class ("`likely match`") and an equal number of majority class instances. Several established classification algorithms are used for WEP, CEP, and CNP, with all of them exhibiting robust performance with respect to their internal configuration.

BLOSS [15] restricts the labeling cost of Supervised Meta-blocking by carefully selecting a training set that is up to 40 times smaller, but retains the original performance. Using *ECBS* weights, it partitions the unlabeled instances into similarity levels and applies rule-based active sampling inside every level. Then, it cleans the sample from nonmatching outliers with high *JS* weights.

Parallelization Approaches. Meta-blocking is adapted to the MapReduce framework in three ways [45]: (1) The *edge-based strategy* stores the blocking graph on the disk, bearing a significant I/O cost. (2) The *comparison-based strategy* builds the blocking graph *implicitly*. A preprocessing job enriches every block with the list of block ids associated with every entity. The Map phase of the second job computes the edge weights and discards all redundant comparisons, while the ensuing Reduce phase prunes superfluous comparisons. This strategy maximizes the efficiency of WEP and CEP and is adapted to Apache Spark in [9]. (3) The *entity-based strategy* aggregates for every entity the bag of all entities that co-occur with it in at least one block. Then, it estimates the edge weight that corresponds to each neighbor based on its frequency in the co-occurrence bag. This approach offers the best implementation for WNP and CNP and their variations (e.g., BLAST). It is adapted to Apache Spark in [151], leveraging the broadcast join for higher efficiency.

To avoid the underutilization of the available resources, these strategies employ *MaxBlock* [45] for load balancing. Based on the highly skewed distribution of block sizes in redundancy-positive block collections, it splits the input blocks into partitions of equivalent computational cost, which is equal to the total number of comparisons in the largest input block.
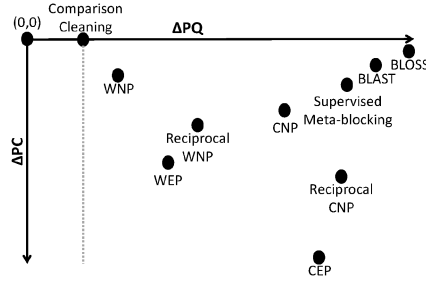
Fig. 6. The relative performance of the main Comparison Cleaning methods.

The *multicore parallelization* of Meta-blocking is examined in [113]. The input is transformed into an array of chunks, with an index indicating the next chunk to be processed. Following the established fork-join model, every thread retrieves the current value of the index and is assigned to process the corresponding chunk. Depending on the definition of chunks, three alternative strategies are proposed: (1) *Naive Parallelization* treats every entity as a separate chunk, ordering all entities in decreasing computational cost (i.e., the aggregate number of comparisons in the associated blocks); (2) *Partition Parallelization* uses MaxBlock to group the input entities into an arbitrary number of disjoint clusters with identical computational cost; and (3) *Segment Parallelization* sets the number of clusters equal to the number of available cores.

4.3 Discussion and Experimental Results

The core characteristic of Block Processing methods is their schema-agnostic functionality, which typically relies on block features, such as size, cardinality, and overlap. This is no surprise, as they are primarily crafted for boosting the performance of schema-agnostic Block Building methods. In fact, extensive experiments demonstrate that Block Processing is indispensable for these methods, raising precision by whole orders of magnitude, at a minor cost in recall [120, 126, 128].

Regarding their relative performance, there is no clear winner among the Block Cleaning methods. For example, both Block Filtering and Block Purging boost *PQ* and *RR* by orders of magnitude, while exhibiting a low computational cost and a negligible impact on *PC* [120, 127]. However, the top performer among them depends not only on their parameter configuration but also on the data at hand [128]. Most importantly, though, Block Cleaning techniques are usually complementary in the sense that multiple ones can be applied consecutively in a single blocking workflow, as depicted in Figure 1(b). For example, Block Filtering is typically applied after Block Purging by lowering *r* to 50% instead of 80%, which is the best configuration when applied independently [126–128].

In contrast, Comparison Cleaning methods are incompatible with each other in the sense that at most one of them can be part of a blocking workflow. The reason is that applying any Comparison Cleaning technique to a redundancy-positive block collection deprives it from its co-occurrence patterns and renders all other techniques inapplicable. These techniques also involve a much higher computational cost than Block Cleaning methods, due to their finer level of granularity. Their relative performance is summarized in Figure 6, based on empirical evidence from experimental studies [128] and individual publications [15, 125–127, 150, 151]. Note that we exclude methods not compared to other Comparison Cleaning techniques (e.g., the techniques presented in [102]).

In more detail, Figure 6 maps the performance of the main Comparison Cleaning methods to a two-dimensional space defined by $\Delta PC = PC(\mathcal{B}') - PC(\mathcal{B})$ on the vertical axis and

$\Delta PQ = PQ(\mathcal{B}') - PQ(\mathcal{B})$ on the horizontal axis, where B and B' stand for the input and the output block collections, respectively. Given that Comparison Cleaning techniques trade lower recall (PC) for higher precision (PQ), ΔPC and ΔPQ take exclusively negative and positive values, respectively. Therefore, the higher a method is placed, the better recall it achieves, whereas the further to the right it lies, the better is its precision. This means that the ideal overall performance corresponds to the upper right corner.

We observe that ΔPC is delimited by two extremes: Comparison Cleaning on the top left corner and CEP on the bottom right corner. The former has no impact on recall, as it increases precision only by removing redundant comparisons. All other Comparison Cleaning techniques discard superfluous comparisons too, thus achieving larger ΔPQ at the cost of a negative ΔPC. On the other extreme, CEP prunes a large portion of superfluous comparisons, yielding very high precision, but the lowest recall. WEP replaces CEP's cardinality constraint with a weight threshold, dropping precision to a large extent for a significantly higher recall. Still, WEP's performance is a major improvement over the input block collection, while being rather robust across numerous datasets. WNP moves further toward this direction, shrinking the decrease in recall and the increase in precision. This is further improved by *Reciprocal WNP*, which significantly raises WNP's precision for slightly lower recall. Thus, it dominates WEP, albeit being sensitive to the characteristics of the data at hand. Compared to CEP, CNP confines its pruning inside individual node neighborhoods. In this way, it achieves a much higher recall for a limited decrease in precision. This is further improved by *Reciprocal CNP*, which reduces CNP's recall slightly for much higher precision and, thus, it often dominates CEP. WNP, CNP, and their variants are improved by *Supervised Meta-blocking* and BLAST, which achieve comparable recall for significantly higher precision. BLAST takes a lead in precision, partially because it employs the most effective weighting scheme, namely Pearson's χ^2 test. Another advantage is that BLAST requires no labeling effort, due to its unsupervised functionality. BLOSS, however, achieves almost perfect recall ($\Delta PC \approx 0$) for the highest precision among all Comparison Cleaning techniques, while requiring merely ~50 labeled instances. Note that exceptions to these general patterns of performance are possible for a particular dataset.

5 FILTERING

Given specific similarity predicates, composing a similarity measure and a corresponding threshold, Filtering techniques receive as input an entity or a block collection and produce as output pairs of entities satisfying these predicates. Next, we present the main filtering methods in the literature, organized in four groups: basic filters proposed by earlier works; prefix filtering and its extensions; partition-based filtering; and methods using tree indexes. An overview of the discussed methods is presented in Table 2, characterized by the type of operation they perform (e.g., search or join), the similarity measure they assume (e.g., token or character based), the type of filters they use (e.g., prefix or partition based), and the index structure they employ (e.g., inverted index or tree).

Basic Filtering. GramCount [57] focuses on incorporating string similarity joins inside a DBMS based on q-grams and edit distance. It is the first work to propose the following techniques:

Length filtering states that if two strings r and s are within edit distance θ, their lengths cannot differ by more than θ. In the case of set similarity joins, the length filter has been adapted to deal with set sizes [7]; e.g., for Jaccard similarity threshold θ, the condition becomes $\theta \cdot |s| \le |r| \le |s|/\theta$. Length filtering is a simple but effective criterion that is employed by many other works alongside more advanced filters. A *position-enhanced* length filter offers a tighter upper bound [94].

Count filtering states that if two strings r and s are within edit distance θ, they must have at least $max(|r|, |s|) - 1 - (\theta - 1) \cdot q$ common q-grams. This filter has also been adapted to sets, in particular in MergeOpt [144], which proposed various optimizations for applying count filtering

Table 2. Overview of String and Set Similarity Join Methods

Method	Operation	Similarity	Filters	Index
GramCount [57]	string join	Edit Distance	length, count, position	q-grams table
MergeOpt [144]	set join	Overlap	count	inverted index
FastSS [18]	string join	Edit Distance	deletion neighborhood	dictionary
SSJoin [22]	set join	Overlap	prefix	DBMS
All-Pairs [12]	vector join	Cosine	prefix	inverted index
DivideSkip [83]	string search	Edit Distance, Overlap	length, position, prefix	inverted index
Ed-Join [180]	string join	Edit Distance	prefix+mismatching q-grams	inverted index
QChunk [133]	string join	Edit Distance	prefix+q-chunks	inverted index
VChunkJoin [176]	string join	Edit Distance	prefix+chunks	inverted index
PPJoin [182, 183]	set join	Overlap	prefix, positional	inverted index
PPJoin+ [182, 183]	set join	Overlap	prefix, positional, suffix	inverted index
MPJoin [140]	set join	Overlap	min-prefix	inverted index
GroupJoin [19]	set join	Overlap	prefix+grouping	inverted index
AdaptJoin [169]	set join	Overlap	adaptive prefix	inverted index
SKJ [177]	set join	Overlap	prefix-based+set relations	inverted index
TopkJoin [181]	top-k set join	Overlap	prefix-based	inverted index
JOSIE [195]	top-k set search	Overlap	prefix, position	inverted index
PartEnum [7]	set join	Hamming, Jaccard	partition based	clustered index
PassJoin [84]	string join	Edit Distance	partition based	inverted index
PTJ [37]	set join	Overlap	partition based	inverted index
B^{ed}−Tree [194]	string search/join	Edit Distance	string orders	B^+-tree
PBI [89]	string search	Edit Distance	reference strings	B^+-tree
MultiTree [192]	set search	Jaccard	tree traversal	B^+-tree
Trie-Join [168]	string join	Edit Distance	subtrie pruning	trie
HSTree [188]	string search	Edit Distance	partition based	segment tree
Trans [193]	top-k set search	Jaccard	transformation distance	R-tree
(a) Exact, centralized, single-predicate algorithms				
FuzzyJoin [2]	set/string join	Hamming, ED, Jaccard	ball-hashing, splitting, anchor points	lookup tables
VernicaJoin [166]	set join	Overlap	prefix, positional, suffix	inverted index
MGJoin [143]	set join	Overlap	multiple prefix	inverted index
MRGroupJoin [37]	set join	Overlap	partition based	inverted index
FS-Join [142]	set join	Overlap	segment based	inverted index
Dima [160, 161]	search, join, top-k	Jaccard, ED	segment based	global & local
(b) Parallel & distributed algorithms				
ATLAS [190]	vector join	Jaccard, Cosine	random permutations	inverted index
BayesLSH [147]	set join	Jaccard, Cosine	All-Pairs/LSH	All-Pairs/LSH
CPSJoin [28]	set join	Jaccard	LSH based	sketches
(c) Approximate algorithms				
LS-Join [173]	local string join	Edit Distance	length, count	inverted index
pkwise [174]	local set join	Overlap	k-wise signatures	inverted index
pkduck [162]	abbreviation matching	Custom	extension of prefix filter	trie
Fast-Join [170]	fuzzy set join	Bipart. graph matching	token-sensitive signatures	inverted index
SilkMoth [36]	fuzzy set join	Bipart. graph matching	weighted token signatures	inverted index
MF-Join [171]	fuzzy set join	Bipart. graph matching	partition based	inverted index
MultiAttr [85]	set search/join	Overlap	tree traversal	prefix tree
Smurf [159]	string matching	Jaccard, Edit Distance	random forest	inverted indexes
AU-Join [184]	string join	Syntactic, Synonym, Taxonomy	pebbles	inverted indexes
(d) Algorithms for complex matching				

with both character-based and token-based similarity measures, and in `DivideSkip` [83], which proposed efficient techniques for merging the inverted lists of signatures.

Position filtering also considers the positions of q-grams in the strings. It states that if two strings r and s are within edit distance θ, a positional q-gram in one cannot correspond to a positional q-gram in the other that differs from it by more than θ positions.

In another line of research, `FastSS` [18] introduces the concept of *deletion neighborhood*, a filtering criterion specifically tailored to edit distance. For a string s, its deletion neighborhood contains substrings of s derived by deleting a certain number of characters. These are then used as signatures for filtering. However, this method is practical only for very short strings.

Prefix-Based Filtering. *Prefix filtering* has been proposed by `SSJoin` [22], which focuses on similarity joins inside a DBMS, and `All-Pairs` [12], which is a main memory algorithm. The prefix filter applies to sets and can also be used for strings represented as sets of q-grams. The elements of each set are first sorted in a global order, typically in increasing order of frequency. Then, the π-prefix of each set is formed by selecting its π first elements in that order. Prefix filter states that for two sets to be similar, their prefixes must contain at least one common element. The prefix size π of a set r is determined based on the similarity measure and threshold being used; e.g., for edit distance threshold θ, $\pi = q \cdot \theta + 1$, while for Jaccard similarity threshold θ, $\pi = \lfloor (1 - \theta) \cdot |r| \rfloor + 1$. As described next, numerous subsequent algorithms have adopted prefix filtering and proposed various optimizations and extensions over it, both for edit distance and set-based similarity joins.

For edit distance, `DivideSkip` [83] uses prefix filtering in combination with length and position filtering, taking special care to efficiently merge the inverted lists of signatures. `Ed-Join` [180] proposes two optimizations based on analyzing the locations and contents of mismatching q-grams to further reduce the prefix length by removing unnecessary elements. `QChunk` [133] introduces the concept of *q-chunks*, which are substrings of length q that start at $1 + i \cdot q$ positions in the string, for $i \in [0, (|r| - 1)/q]$. Given two strings r and s, QChunk extracts q-grams from the one and q-chunks from the other; if r and s are within edit distance θ, the size of the intersection between the q-grams of r and the q-chunks of s should be at least $\lceil |s|/q \rceil - \theta$. `VChunkJoin` [176] uses nonoverlapping substrings called *chunks*, ensuring that each edit operation destroys at most two chunks. This yields a tight lower bound on the number of common chunks that two strings must share if they match.

For set similarity joins, `PPJoin` [182, 183] extends prefix filtering with *positional filtering*. This takes also into consideration the positions where the common tokens in the prefix occur, thus deriving a tighter upper bound for the overlap between the two sets. In addition, `PPJoin+` [182, 183] further uses *suffix filtering*. Following a divide-and-conquer strategy, this partitions the suffix of the one set into two subsets of similar sizes. The token separating the two partitions is called *pivot* and is used to split the suffix of the other set. This allows one to calculate the maximum number of tokens in each pair of corresponding partitions between the two sets that can match.

`MPJoin` [140] adds a further optimization over `PPJoin` that allows for dynamically pruning the length of the inverted lists. This reduces the computational cost of candidate generation, rather than the number of candidates. `GroupJoin` [19] extended `PPJoin` with *group filtering*, whose candidate generation treats all sets with identical prefixes as a single set. Multiple candidates may thus be pruned in batches. `AdaptJoin` [169] proposed *adaptive prefix filtering*, which generalizes prefix filtering by adaptively selecting an appropriate prefix length for each set. It supports longer prefixes dynamically, extending their length by $n - 1$, and then prunes a pair of sets if they contain less than n common tokens in their extended prefixes. Prefix filtering is a special case where $n = 1$.

A different perspective for speeding up set similarity joins is proposed by `SKJ` [177]. The idea is based on the following observation: existing approaches examine each set individually when computing the join; however, it is possible to improve efficiency through computational cost sharing

between *related sets*. To this end, the SKJ algorithm introduces *index-level skipping*, which groups related sets in the index into blocks, and *answer-level skipping*, which incrementally generates the answer of one set from an already computed answer of another related set.

Finally, there are Filtering techniques for computing top-k results progressively, instead of requiring the user to select a similarity threshold. TopkJoin [181] retrieves the top-k pairs of sets ranked by their similarity score, based on prefix filtering and on the monotonicity of maximum possible scores of unseen pairs. JOSIE [195] presents a method for top-k set similarity search. It exploits prefix and position filtering, but instead of dealing with sets of relatively small size (e.g., ~100 tokens), it is crafted for finding joinable tables in data lakes, where sets represent the distinct values of a table column, comprising millions of tokens. This introduces new challenges, which are tackled by proposing an algorithm that minimizes the cost of set reads and inverted index probes.

Partition-Based Filtering. The algorithms in this category partition each string or set into multiple disjoint segments in such a way that matching pairs have at least one common segment. PartEnum [7] generates a signature scheme based on the principles of *partitioning* and *enumeration*. The former states that if two vectors with Hamming distance not higher than k are partitioned into $k + 1$ equi-sized partitions, then they must have at least one common partition. The latter states that if these vectors are partitioned instead into $n > k$ equi-sized partitions, then they must have in common at least $n - k$ partitions. PassJoin [84] partitions a string into a set of segments and creates inverted indices for the segments. Then, for each string, it selects some of its substrings and uses them to retrieve candidates from the index. A method is proposed to minimize the number of segments required to find the candidate pairs. PTJ [37] proposes an approach to increase the pruning power of partition-based filtering by using a mixture of the subsets and their one-deletion neighborhoods, which are subsets derived from a set after eliminating one element.

Essentially, these methods are based on the *pigeonhole principle*, which states that if n items are contained in m boxes, at least one box has no more than $\lfloor n/m \rfloor$ items. This is extended by the *pigeonring principle* [134], which organizes the boxes in a ring and constrains the number of items in multiple boxes rather than a single one, thus offering tighter bounds. Applying it to various similarity search problems shows that pigeonring always produces fewer or an equal number of candidates than the pigeonhole principle does and that pigeonring-based algorithms can be implemented on top of existing pigeonhole-based ones with minor modifications [134].

Tree-Based Filtering. Most methods presented so far build inverted indexes on the signatures extracted from the strings or sets. Next, we present algorithms employing tree-based indexes.

Most approaches are based on the B$^+$-tree. B^{ed}-Tree [194] proposes a B$^+$-tree-based index for range and top-k similarity queries as well as similarity joins, using edit distance. It is based on a mapping from the string space to the integer space to support efficient searching and pruning. PBI [89] uses a B$^+$-tree index and exploits the fact that edit distance is a metric. The string collection is partitioned according to a set of selected *reference strings*. Then, the strings in each partition are indexed based on their distances to their corresponding reference strings. The proposed approach supports both range and k-NN queries and can be integrated inside a DBMS. In MultiTree [192], each element in a set is represented as a vector and is mapped to an integer according to a defined global ordering, which is then used to insert the element in the B$^+$-tree index. Searching for similar elements is then done via a range query on the index.

In another line of research, Trie-Join [168] proposes a trie-based technique for string similarity joins with edit distance. Each trie node represents a character in the string. Thus, strings with a common prefix share the same ancestors. A trie node is called an *active node* of a string s if their edit distance is not larger than the given threshold. This leads to a technique called *subtrie pruning*: given a trie T and a string s, if node n is not an active node for every prefix of s, then n's descendants cannot be similar to s. HSTree [188] recursively partitions strings into disjoint segments and builds

a hierarchical segment tree index. This is then used to support both threshold-based and top-k string similarity search based on edit distance. Finally, a transformation-based framework for top-k set similarity search is presented in [193]. It transforms sets of various lengths into fixed-length vectors in such a way that similar sets are mapped closer to each other. An R-tree is then used to index these records and prune the space during search.

5.1 Parallel and Distributed Algorithms

MapReduce-based approaches have been proposed to tackle scalability issues when dealing with very large collections of sets or strings. A theoretical analysis of different methods for performing similarity joins on MapReduce is presented in [2]. It considers algorithms that operate in a single MapReduce job, avoiding the overhead associated with initiating multiple ones. It shows that different algorithms provide different tradeoffs with respect to map, reduce, and communication cost.

VernicaJoin [166] is based on prefix filtering. It computes prefix tokens and builds an inverted index on them. Then, it generates candidate pairs from the inverted lists, using additionally the length, positional, and suffix filters to prune candidates. A deduplication step is finally employed to remove duplicate result pairs generated from different reducers. MGJoin [143] follows a similar approach to VernicaJoin but introduces multiple prefix orders and a load balancing technique that partitions sets based on their length. MRGroupJoin [37] is a MapReduce extension of PTJ [37]. It applies a partition-based technique, where records are grouped by length and are partitioned in subrecords, such that matching records share at least one subrecord. The process is performed in a single MapReduce job. FS-Join [142] sorts the tokens in each set in increasing order of frequency, and then splits each set into disjoint subsets using appropriate pivot tokens. These subsets are then grouped together so that subsets from different groups are nonoverlapping.

Finally, Dima [160, 161] is a distributed in-memory system built on top of Spark that supports threshold and top-k similarity search and join with both token-based and character-based similarities. It relies on signature-based global and local indexes for efficiency. The proposed signatures are adaptively selectable based on the workload, which allows one to balance the workload among partitions. Dima extends the Catalyst optimizer of Spark SQL to introduce cost-based optimizations.

5.2 Approximate Algorithms

Approximate algorithms for similarity search and join increase the efficiency of the Filtering step at the cost of allowing both false positives and false negatives, thus missing some matches [10, 172]. They typically rely on *locality-sensitive hashing* (LSH) [55], which transforms an item to a low-dimensional representation such that similar items have much higher probability to be mapped to the same hash code than dissimilar ones. This property allows LSH to be exploited in the filtering phase to generate candidates [32, 90, 163]. The basic idea is that each object is hashed several times using randomly chosen hash functions. Then, candidates are those pairs of objects that have been hashed to the same code by at least one hash function.

ATLAS [190] is a probabilistic algorithm that is based on random permutations both to generate candidates and to estimate the similarity between candidate pairs. It also proposes a method to efficiently detect cluster structures within the data, which are then exploited to search for similar pairs only within each cluster. BayesLSH [147] combines Bayesian inference with LSH to estimate similarities to a user-specified level of accuracy. It uses LSH for both Filtering and Verification, providing probabilistic guarantees on the resulting accuracy and recall. CPSJoin [28] is a randomized algorithm for set similarity joins. It uses a recursive filtering technique, building upon a previously proposed index for set similarity search [27], as well as sketches for estimating set similarity. The algorithm has 100% precision and provides a probabilistic guarantee on recall.

5.3 Algorithms for Complex Matching

The works discussed so far assume a single similarity predicate; i.e., they apply to the values of a specific attribute. Moreover, when comparing sets, they assume binary matching between their elements, while in the case of strings, they compare strings in their entirety. In the following, we present methods that employ *multiple* similarity predicates or more complex ones.

Local Matching. A local string similarity join finds pairs of strings that contain similar *substrings*. Under edit distance constraints, it can be defined as matching any l-length substring with up to k errors. LS-Join [173] is based on the observation that if two strings are locally similar, they must share at least one common q-gram, for a suitably calculated gram length q. An inverted index is constructed incrementally during the search. For every examined string, its q-grams are generated and the candidates are retrieved from the index by finding those strings that have matching q-grams.

pkwise [174] detects pieces of text in a given collection that share similar *sliding windows*, i.e., multisets containing w consecutive tokens of a given document. The similarity of two sliding windows is defined as the overlap of those sets. Prefix filtering is used, but instead of relying on single tokens to build the signatures, it proposes *k-wise signatures*, which comprise combinations of k tokens. Larger values of k increase the signatures' selectivity but also the cost of signature generation. An additional optimization is to share common signatures across adjacent windows.

Finally, pkduck [162] matches strings with *abbreviations*, based on a new similarity measure that accounts for abbreviations. It also proposes an appropriate signature scheme that extends prefix filtering and generates signatures without iterating over all strings derived from an abbreviation.

Fuzzy Matching. Rather than assuming a binary match, in this setting, the similarity between the elements of two sets may take any value between 0 and 1. In fact, it is defined as the maximum matching score in the bipartite graph representing the matches between their elements.

In Fast-Join [170], edge weights in this bipartite graph denote the edit similarities between matching elements. The proposed method follows the filter verification framework, creating a signature for each set such that matching sets have overlapping signatures. The signature of a set is composed of an appropriately selected subset of its tokens. SilkMoth [36] generalizes and improves upon this work, providing a formal characterization of the space of valid signatures. Given that finding the optimal signature is NP-complete, it proposes heuristics to select signatures. To further reduce candidates, a refinement step is added: it compares each set with its candidates and rejects those for which certain bounds do not hold. Both edit distance and Jaccard coefficient are supported for measuring the similarity between elements. MF-Join [171] performs element- and record-level filtering. The former utilizes a partition-based signature scheme with a frequency-aware partition strategy, while the latter exploits count filtering and an upper bound on record-level similarity.

Multiple Predicates. A method for similarity search and join on *multiattribute* data is presented in [85]. For instance, given an entity collection where each entity is described by its name and address, this work identifies pairs of entities having *both* similar names and similar addresses. To enable simultaneous filtering on multiple attributes, a combined prefix tree index is built on these attributes. The construction of the index is guided by a cost model and a greedy algorithm. In another direction, Smurf [159] performs string matching between two collections of strings based on multiple-predicate matching conditions in the form of a *random forest* classifier that is learned via active learning. Filtering techniques for string similarity joins are exploited to speed up the execution of the random forest. The focus and novelty of this work is on how to reuse computations across the trees in the forest to further increase efficiency. Finally, AU-Join [184] presents a new framework for string similarity joins that supports not only syntactic similarity measures,

such as Jaccard similarity on *q*-grams, but also *semantic* similarities, including *synonym-based* and *taxonomy-based* matching. It partitions strings into segments and applies different types of similarity measures on different pairs of segments. A new signature scheme, called *pebble*, handles multiple similarity measures: pebbles are *q*-grams for gram-based similarity, the left-hand side of a synonym rule for synonym similarity, and ancestor nodes in the taxonomy for taxonomy similarity.

5.4   Discussion and Experimental Results

Filtering techniques for string and set similarity joins have attracted a lot of research interest over the past two decades. Early works view this operation as an extension of the standard join operator in relational databases, where the join condition is based on similarity rather than equality [22, 57]. The same perspective is shared by more recent works, like those proposing B$^+$-tree-based indexes, which can be easily integrated into an existing DBMS [89, 192, 194]. Another characteristic example is Dima [160, 161], which extends the Catalyst optimizer of Spark SQL to support similarity-based queries. In this sense, similarity joins are sometimes referred to as *approximate* or *fuzzy* joins, although this should not be confused with the approximate algorithms in Section 5.2 or the fuzzy set joins in Section 5.3. Numerous Filtering techniques have been proposed by more recent works, which focus on main memory execution. *Prefix-based* filtering is the most popular approach [12, 19, 22, 83, 133, 140, 169, 176, 180, 182, 183], followed by *partition-based* filtering [7, 37, 84]. Furthermore, to scale similarity joins to large collections, distributed [2, 37, 142, 143, 160, 161, 166] and approximate [28, 147, 190] algorithms have been proposed.

More recently, there has been increasing focus and interest on works that deal with more complex similarity predicates. These include the matching of strings based on *substrings* or *abbreviations* [162, 173, 174], matching of sets based on *fuzzy matching* of their elements [36, 170, 171], and the combination of *multiple* similarity predicates [85, 159]. These works can be considered as more closely relevant to matching entity profiles in Entity Resolution.

Regarding performance, a series of experimental analyses provides interesting insights [50, 64, 95]. However, each study focuses on a certain subset of the aforementioned methods. Below, we briefly summarize their findings, including additional results from individual papers to fill the gaps.

Similarity joins using Edit Distance. A comparison between `FastSS`, `All-Pairs`, `DivideSkip`, `Ed-Join`, `QChunk`, `VChunkJoin`, `PPJoin`, `PPJoin+`, `AdaptJoin`, `PartEnum`, `PassJoin`, and `Trie-Join` is conducted in [64]. The results demonstrate that `PassJoin` is the most efficient algorithm, with `FastSS` providing a reliable alternative in the case of very short strings.

Similarity joins using set-based measures. `AdaptJoin` and `PPJoin+` are reported as the best algorithms in the aforementioned study [64]. Different results, though, are reported in a subsequent study that compares `All-Pairs`, `PPJoin`, `PPJoin+`, `MPJoin`, `AdaptJoin`, and `GroupJoin`. It indicates that the plain prefix filtering, i.e., `All-Pairs`, is still quite competitive, winning in the majority of cases. `PPJoin` and `GroupJoin` exhibit the best median and average performance, respectively, while more sophisticated filters are found to provide only moderate improvements in some cases or even to negatively affect performance. The difference with the results in [64] is attributed to the more efficient verification step; reducing the cost of Verification means that complex and, thus, time-consuming filters often do not pay off, despite reducing the number of candidate pairs.

Prefix vs. partition filtering. PTJ is compared against `PPJoin+` and `AdaptJoin` in [37], showing that it outperforms both methods. The same comparison is performed in [177], showing that PTJ does not outperform those methods in most cases. As noted in [177], this discrepancy seems to be caused by differences in implementation; specifically, the comparison in [37] uses the original implementations of `PPJoin+` and `AdaptJoin`, while the one in [177] uses the optimized implementations provided by [95]. Overall, PTJ may generate fewer candidates but uses complex index

Fig. 7. (a) The taxonomy of the hybrid, join-based blocking methods. (b) Timeline of the landmarks in the evolution of Blocking, Filtering, and their convergence.

structures, thus spending much more time on the filtering phase compared to prefix-based algorithms. Another factor that affects the performance of prefix filtering is the frequency distribution of the tokens in the dataset. The core idea of prefix filtering is to select rare tokens as signatures so as to reduce the number of candidates. In [64], an experiment involving different dataset distributions shows that PPJoin(+) and AdaptJoin perform better in datasets with Zipfian distribution than a uniform one.

Set relations. Another interesting finding is that set relations can be effectively exploited to speed up the computation of similarity joins [177]. In the presented experiments, the proposed algorithm, SKJ, consistently outperforms PPJoin, PPJoin+, AdaptJoin, and PTJ across all datasets.

Tree-based algorithms. These algorithms typically focus on similarity search rather than join. HSTree and PBI are compared against B^{ed}-Tree in [188] and [89], respectively, reporting better performance. Also, Trans shows better performance than MultiTree in [193]. BiTrieJoin, an improved variant of TrieJoin, is reported in [64] to have comparable performance to PassJoin for short strings, but it underperforms for medium and long strings.

Distributed algorithms. VernicaJoin, MGJoin, MRGroupJoin, and FS-Join are experimentally compared in [50]. VernicaJoin exhibits the best performance in most cases, but all algorithms are often outperformed by nondistributed ones. This should be attributed to the overhead introduced by the MapReduce framework as well as to high or skewed data replication between map and reduce tasks. The latter constitutes an inherent limitation of the distributed algorithms that cannot be overcome by simply increasing the number of nodes in the cluster. In [160], Dima is shown to outperform the adaptation of VernicaJoin to Apache Spark.

Approximate algorithms. The experimental survey in [64] included a comparison between BayesLSH-lite and exact algorithms. Moreover, ATLAS, BayesLSH, and CPSJoin have been compared against All-Pairs in [190], [147], and [28], respectively. Overall, the experiments indicate that approximate algorithms are preferable for low-similarity thresholds, e.g., for Jaccard similarity below 0.5, while exact algorithms perform better for high thresholds.

6 JOIN-BASED BLOCKING METHODS

We now elaborate on Block Building methods that incorporate Filtering techniques, converting Blocking into a nearest neighbor search. As illustrated in Figure 7(a), we categorize these hybrid techniques into three major categories according to the filtering techniques they employ: the *lossless* ones rely on exact, single-predicate filtering techniques (cf. Table 2(a)); the *lossy* ones rely on approximate filtering (cf. Section 5.2); and the *spatial* ones leverage spatial join techniques for filtering. Note that the lossy hybrid methods are further distinguished into *static* and *dynamic* ones, depending on whether they are independent or interwoven with Matching, respectively.

Starting with the lossless hybrid methods, the simplest approach is to combine Prefix Filtering with Token Blocking, creating one block for every token that appears in the prefix of at least two

entities [29]. Another approach is *Adaptive Filtering* [59], which couples schema-aware, nonlearning Block Building techniques with two filtering methods. First, blocks are created by extracting keys from specific attributes. In every block with a size exceeding a predetermined threshold, Length and Count Filtering are applied for Comparison Cleaning, using an edit distance threshold on an attribute that is not considered by the initial transformation function.

Another lossless hybrid method is *LIMES*, which operates only on metric spaces [107]. Its core idea is to leverage the triangle inequality to approximate the distance between entities based on previous comparisons. Utilizing sets of entities as reference points, called *exemplars*, this method computes lower and upper bounds to filter out superfluous comparisons before their execution.

In another direction, *MultiBlock* [62] optimizes the execution of complex matching rules that are composed of special similarity functions for textual, geographic and numeric values. A block collection is created for every similarity function such that similar entities share multiple blocks; e.g., edit distance is supported for textual values and blocks are created for character q-grams such that entity pairs satisfying the distance threshold co-occur in a sufficient number of blocks. Then, all block collections are aggregated into a multidimensional index that respects the co-occurrence patterns of similar entities and guarantees no false dismissals, i.e., $PC = 1$.

Regarding the lossy approaches, they are dominated by techniques based on LSH [55], which efficiently estimates the similarity between two attribute values v_i and v_j by randomly sampling hash functions f from a *sim-sensitive* family F such that the probability $Pr(f(v_i) = f(v_j))$ equals $sim(v_i, v_j)$ for any pair of attribute values and any function $f \in F$. This means that LSH derives $sim(v_i, v_j)$ from the proportion of hash functions f such that $f(v_i) = f(v_j)$. Typically, the required number of these functions is relatively small for a sufficiently small sampling error; e.g., for 500 functions, the maximum sampling error is about ±4.5% with 95% confidence interval [42].

In the context of ER, LSH is typically combined with MinHash signatures [20], which efficiently estimate the Jaccard similarity as follows [58, 158]. Given an entity collection $\mathcal{E}$, the values of selected attribute names are converted into a bag of *k-shingles*, i.e., k consecutive words or characters. Then, a matrix M of size $K \times |\mathcal{E}|$ is formed, with the rows corresponding to the K distinct shingles that appear in all attribute values and the columns to the input entities. The value of every cell $M(i, j)$ indicates whether the entity e_j contains the shingle s_j, $M(i, j) = 1$, or not, $M(i, j) = 0$. Given that M is a sparse matrix, p random minhash functions are used to reduce its dimensionality: they are applied to each column, deriving a new matrix M' of size $p \times |\mathcal{E}|$. The p rows are then partitioned into b nonoverlapping bands and a hash function is applied to every band of each column. The resulting buckets are treated as blocks that provide probabilistic guarantees that the pairs of similar entities co-occur in at least one block. In fact, the desired probabilistic guarantees can be used for configuring the parameters of LSH, i.e., the number of hash functions, rows, and bands [58].

In this context, LSH is combined with K-Means in KLSH [158]. K-Means is applied to the low-dimensional columns of M', which represent the input entities. The resulting clusters form a disjoint block collection $\mathcal{B}$, with $|\mathcal{B}|$ determined by the desired average number of entities per block.

LSH is applied to the distributed representations (i.e., embeddings) of the input entities in *DeepER* [43]. Every entity is transformed into a dense, real-valued vector by aggregating the embeddings of all attribute value tokens, which are pretrained by word2vec [101], Glove [131], and so forth. This vector is then hashed into multiple buckets with LSH. A block is then created for every entity containing its top-N most likely matches, which are detected using Multiprobe-LSH [90].

LSH is also combined with a semantic similarity in SA-LSH (i.e., semantic-aware LSH) [175]. A taxonomy tree is used to model the concepts that describe the input entity collection. The semantic similarity of two entities is inversely proportional to the length of the paths that connect the

corresponding concepts and their children: the longer the paths, the lower the semantic similarity. The concepts of every entity are converted into a hash signature through a semantic hashing algorithm. The resulting low-dimensional signatures are directly combined with the signatures that are extracted from the n-grams of selected attribute values, capturing the textual similarity of entities. However, the construction of the taxonomy tree requires heavy human intervention.

Regarding the dynamic lossy methods, LSH is combined with R-Swoosh [14] in [93] through a MapReduce parallelization. Initially, a job is used for defining blocks using LSH. Then, a graph-parallel Pregel-based platform applies R-Swoosh, iteratively executing the nonredundant comparisons in the blocks and computing the transitive closure of the detected duplicates.

LSH also lies at the core of cBV-HB [70, 71], which embeds the textual values of selected attributes into a compact binary Hamming space that is efficient, due to the limited size of its embeddings (e.g., 120 bits for four attributes), and preserves the original distances in the sense that certain types of errors correspond to specific distance bounds. Special care is taken to support composite matching rules that involve the main logical operators (i.e., AND, OR, and NOT).

Similarly, *HARRA* [76] uses LSH to hash similar entities into the same buckets. Inside every bucket, all pairwise comparisons are executed and duplicates are merged into new profiles. The new profiles are hashed into the existing hash tables and the process is repeated until no entities are merged or another stopping criterion is met (e.g., the portion of merged profiles drops below a predetermined threshold). In every iteration, special care is taken to avoid redundant and superfluous comparisons.

Finally, spatial hybrid methods combine spatial joins with Block Building. The core approach is *StringMap* [65], which converts schema-aware blocking keys to a similarity-preserving Euclidean space, whose dimensionality d is heuristically derived from a random sample (typically, $d \in [15, 25]$). For each dimension, a linear algorithm initially selects two pivot attribute values that are (ideally) as far apart as possible. Subsequently, the coordinates of all other attribute values are determined through a comparison with the pivot strings. Using an R-tree or a grid-based index in combination with two weight thresholds, similar attribute values are clustered together into overlapping blocks.

This approach is enhanced by *Extended StringMap* [25], which replaces the weight thresholds with cardinality ones, and the *Double embedding scheme* [1]. The latter initially maps the input entities to the same d-dimensional Euclidean space. Next, the embedded attribute values are mapped to another Euclidean space of lower dimensionality $d' < d$. A similarity join is performed in the second Euclidean space using a k-d tree index. The resulting candidate matches are then clustered in the first, d-dimensional Euclidean space. The experimental study suggests that the d'-dimensional space significantly reduces the runtime of StringMap by 30% to 60%.

7 BLOCKING VS. FILTERING: COMMONALITIES AND DIFFERENCES

The timeline in Figure 7(b) summarizes the landmarks in the evolution of the two frameworks showing their gradual convergence. We observe that Blocking is the oldest discipline, with the first relevant technique, namely SB, presented in 1969 [49]. For several decades, research focused on schema-based techniques, with the most significant breakthrough taking place in 1995, with the introduction of SN [60]. The first schema-agnostic Block Building technique is *Semantic Graph Blocking* [109], introduced in 2007, but it considers only entity links. In 2011, it was followed by TB [116], which exclusively applies to textual values. Block Processing was introduced in 2009 by Iterative Blocking [179], followed by the use of Canopy Clustering for Blocking in 2012 [25] and the introduction of Meta-blocking in 2014 [121]. For Filtering, the first similarity join to be used in an RDBMS can be traced back to 2001 [57], while the techniques for in-memory execution were coined in 2007 [12]. Attempts to further increase efficiency by allowing approximate results were

first presented in 2011 [190]. The first works on massive parallelization for Filtering appeared in 2010 [166], for Blocking in 2012 [77], and for Block Processing in 2015 [44]. The convergence of the two frameworks essentially started in 2011 with MultiBlock [62], which introduced Join-based Blocking, whereas multiple-predicate Filtering for efficient Matching was introduced by Smurf in 2018 [159].

Regarding the qualitative comparison of the two frameworks, we observe that they have a number of commonalities: (1) Both serve the same purpose: they increase ER efficiency by reducing the number of performed comparisons. To this end, both employ a stage producing candidate matches, which are subsequently examined analytically in order to remove false positives. (2) Both usually operate either on two clean but overlapping data collections (Record Linkage for Blocking, Cross-table Join for Filtering) or on a single dirty data collection (Deduplication for Blocking, Self-join for Filtering). (3) Both extract signatures such that the similarity of two entities is reflected in the similarity of their signatures. (4) Both also apply similar implementation-level optimizations, representing signatures with integer ids, instead of strings, so as to reduce the memory footprint and facilitate in-memory execution. (5) Both include character- and token-based methods. For Blocking, the former methods mainly pertain to schema-aware techniques that apply character-level transformations to blocking keys (e.g., q-grams, suffixes, etc.), while token-based methods primarily pertain to schema-agnostic methods. For Filtering, similarity measures can also be distinguished between character-based (e.g., edit similarity) and token-based ones (e.g., Jaccard), even though many algorithms can be adapted to handle both. (6) In both cases, textual data have been combined with other types of data, particularly with spatial or spatiotemporal data, including [106] for Blocking and [13, 19, 52, 63] for Filtering. (7) Both can be used in real-time applications, where the input consists of a query entity and the goal is to identify the most similar ones in the minimum possible time. This is called Similarity Search in the case of Filtering and Real-time ER in the case of Blocking (see Section 9 for more details).

Due to these commonalities, several works use the two frameworks interchangeably, considering Filtering as a means for Blocking (e.g., [29]). In reality, though, Blocking and Filtering have several distinguishing characteristics: (1) By definition, a blocking scheme applies to a single entity, considering all its attribute values (schema-agnostic methods), or combinations of multiple values (schema-aware techniques). In contrast, Filtering usually applies to a pair of values from the same attribute of two entities. (2) Blocking relies on positive evidence, clustering together similar entities, while Filtering relies on negative evidence, detecting dissimilar entities early on. (3) Blocking is typically independent of Entity Matching, whereas Filtering is interwoven with it, as its goal is to optimize the execution of a matching rule. (4) Blocking is an inherently approximate procedure that falls short of perfect recall (*PC*), even when providing probabilistic guarantees (e.g., LSH Blocking in DeepER [43]). In contrast, most Filtering methods provide an exact solution, returning all pairs of values that exceed the predetermined threshold along with false positives. (5) Blocking trades slightly lower recall (*PC*) for much higher precision (*PQ*), while Filtering trades filtering power for filtering cost. (6) Blocking may be modeled as a learning problem, where the goal is to define supervised blocking schemes that simultaneously optimize *PC*, *PQ*, and *RR*, but Filtering requires no labeled set for learning to mark a comparison as true negative. Instead, it relies on a theoretical analysis based on the given similarity measure and threshold. (7) Preserving privacy is orthogonal to Filtering, with very few works examining privacy-preserving similarity joins [67, 87, 189]. In contrast, Blocking constitutes an integral part of privacy-preserving ER, with several relevant works (for details, refer to a recent survey [165]). (8) Blocking constitutes an integral part of pay-as-you-go ER applications, conveying a significant body of relevant works, as described below. This does not apply to Filtering, given that the only relevant technique is TopkJoin [181].

Regarding the quantitative comparison between Blocking and Filtering, few works have actually examined their relative performance. The two frameworks are experimentally juxtaposed in [153–155] in terms of effectiveness and time efficiency. Using a series of real-world datasets, RDFKeyLearner is compared against AllPairs, PPJoin(+), and EdJoin in [153, 154] and against EdJoin, PPJoin+, and FastJoin in [155]. All methods are fine-tuned using a sample of each dataset. The outcomes indicate no significant difference in effectiveness, but regarding time efficiency, Filtering is consistently faster in generating candidate matches and consistently slower in executing the corresponding pairwise comparisons, due to their larger number. In [155], the relative scalability of RDFKeyLearner and EdJoin is examined over synthetic datasets of 10^5, $2 \cdot 10^5$, ..., 10^6 entities. Again, EdJoin produces more candidate matches and thus is slower than RDFKeyLearner.

In [82], an experimental analysis over four real-world datasets investigates the combined effect of Blocking and Filtering on ER efficiency, implementing the workflow in Figure 1(b). The results suggest that together, the two frameworks reduce the overall ER running time from 33% to 76%, with an average of 50%. However, only one method per framework is considered: the manually fine-tuned SB and PPJoin in combination with Cosine and Jaccard similarity. Note that due to its careful, manual fine-tuning, Blocking has no impact on ER effectiveness.

However, more experimental analyses are required for drawing safe conclusions about the relative performance of Blocking and Filtering. These analyses should include a large, representative variety of techniques per framework along with several established benchmark datasets and should examine the benefits of combining the two frameworks in more depth.

8 BLOCKING AND FILTERING IN ENTITY RESOLUTION SYSTEMS

We now present the main systems that address ER, examining whether they incorporate any of the aforementioned methods to improve the runtime and the scalability of their workflows. We analytically examined the 18 noncommercial and 15 commercial systems listed in the extended version of [80][1] along with the 10 Link Discovery frameworks surveyed in [105]. Table 3 summarizes the characteristics of 12 open-source ER systems that include at least one Blocking or Filtering method.

Half of the tools offer a graphical user interface and are implemented in Java. Regarding the type of the input data, most systems support structured data. The only exceptions are the three Link Discovery frameworks, which are crafted for semistructured data. JedAI is the only tool that applies uniformly to both structured and semistructured data.

We also observe that all systems include Blocking methods, with Standard Blocking (SB) and Sorted Neighborhood (SN) being the most popular ones. The first four systems are Link Discovery frameworks that implement custom approaches: KnoFuss and SERIMI apply Token Blocking only to the literal values of RDF tiples, while Silk and LIMES implement hybrid methods, MultiBlock and LIMES, respectively (see Section 6). Febrl and JedAI offer the largest variety of established techniques. The former provides their original, schema-aware implementation, while the latter provides their schema-agnostic adaptations. For this reason, JedAI is the only tool that implements Block Processing techniques as well.

Note that Block Building is also a core part of the ER workflow in several commercial systems, such as IBM Infosphere and Informatica Data Quality [80]. These systems are generally required to handle diverse types of data, focusing on data exploration and cleaning. They typically provide variations of SB, allowing users to extract blocking keys from specific attributes through a sophisticated GUI that provides statistics and data analysis. As a result, users' expertise and experience with specific domains is critical for the performance of these systems' blocking components.

---

[1]The extended version of [80] is available at http://pages.cs.wisc.edu/~anhai/papers/magellan-tr.pdf.

Table 3. Blocking and Filtering Methods in Open-Source Systems for Entity Resolution

Tool	Blocking	Filtering	GUI	Language	Data Formats
KnoFuss [108]	Literal Blocking	-	No	Java	RDF, SPARQL
SERIMI [8]	Literal Blocking	-	No	Ruby	SPARQL
Silk [167]	Multiblock	-	Yes	Scala	RDF, SPARQL, CSV
LIMES [107]	Custom methods	PPJoin+, EdJoin, custom methods, e.g., ORCHID [106]	Yes	Java	RDF, SPARQL, CSV
Dedupe [17]	SB with learning-based techniques	-	No	Python	CSV, SQL
DuDe [39]	SB, SN, Sorted Blocks	-	No	Java	CSV, JSON, XML, BibTex, Databases (Oracle, DB2, MySQL and PostgreSQL)
Febrl [23]	SB, SN, Sorted Blocks, Suffix Arrays, Extended Q-Grams, Canopy Clustering, StringMap	-	Yes	Python	CSV, text-based
FRIL [66]	SB, SN	-	Yes	Java	CSV, Excel, COL, Database
OYSTER [104]	SB	-	No	Java	Text-based
RecordLinkage [145]	SB (with SOUNDEX)	-	No	R	Database
Magellan [80]	SB, SN, it also supports user-specified blocking methods	Overlap, Length, Prefix, Position, Suffix	Yes	Python	CSV
JedAI [129]	SB, SN, Extended SN, Suffix Arrays, Extended Suffix Arrays, LSH, Q-Grams, Extended Q-Grams + Block Processing	To be added in the forthcoming version 3	Yes	Java	CSV, RDF, SPARQL, XML, Database

Surprisingly, only two systems currently include Filtering algorithms for improving the runtime of their matching process: LIMES and Magellan. The latter actually offers the largest variety of established techniques through the py_stringsimjoin package. Filtering techniques are also provided by FEVER [81], which is a closed-source ER tool, as well as by JedAI's forthcoming version 3. Still, a mere minority of ER tools enables users to combine the benefits of Blocking and Filtering, despite the promising potential of their synergy (see below for more details). Most importantly, these tools exclusively consider traditional Filtering algorithms that apply to the values of individual attributes. Hence, they disregard the recent Filtering techniques for Complex Matching (cf. Section 5.3), which are more suitable for Entity Resolution. Therefore, more effort should be devoted to developing ER tools that make the most of the synergy between Blocking and Filtering.

9 FUTURE DIRECTIONS

Various directions seem promising for future work, from entity evolution [115] to deep learning [43] and summarization algorithms [69], which minimize the memory footprint of blocks, while accelerating their processing. The following are more mature fields, having assembled a critical mass of methods already.

Progressive Entity Resolution. Due to the constant increase of data volumes, new *progressive* or *pay-as-you-go* ER methods provide the best possible *partial solution* within a limited budget of temporal or computational resources. Based on Blocking, they schedule the processing of entities, comparisons, or blocks according to the likelihood that they involve duplicates.

Among the schema-based methods, *Progressive Sorted Neighborhood* (PSN) [178] applies an incremental window size w to a sorted list of entities until reaching the available budget. *Dynamic PSN* [130] adjusts PSN's processing order on the fly, according to the results of an oracle, while *Hierarchy of Record Partitions* [178] creates a hierarchy of blocks that is resolved level by level,

from the leaves, which contain the most likely matches, to the root. A variation of this approach is adapted to MapReduce in [6], while the *Ordered List of Records* [178] converts it into a list of entities that are sorted by their likelihood to produce matches. A progressive solution for relational Multi-source ER over different entity types is proposed in [5]. P-RDS adapts LSH-based blocking to a progressive operation by rearranging the processing order of hash tables according to the number of matching and unnecessary comparisons in the buckets examined so far.

Among the schema-agnostic methods, *Local Schema-agnostic Progressive SN* [152] slides an incremental window over the sorted list of entities that is created by schema-agnostic SN, and for each window size, it orders the nonredundant comparisons according to the co-occurrence frequency of their entities. *Global Schema-agnostic Progressive SN* [152] does the same for a predetermined range of windows, eliminating all redundant comparisons they contain. *Progressive Block Scheduling* [152] orders the blocks in ascending number of comparisons and then prioritizes all comparisons per block in decreasing weight. *Progressive Profile Scheduling* [152] orders entities in decreasing average comparison weight and then prioritizes all comparisons per entity in decreasing weight.

The schema-agnostic methods excel in recall and precision [152] but exclusively support *static* prioritization, defining an immutable processing order that disregards the detection of duplicates. Hence, more research is needed for developing *dynamic schema-agnostic* progressive methods.

Real-Time Entity Resolution. This task matches a query entity to the available entity collections in (ideally) subsecond runtime. An early solution is presented in [26], which precalculates the similarities between the attribute values of entities co-occurring in SB blocks to avoid similarity calculations at query time. Its indexes are dynamically adapted to query entities in [138].

Other dynamic indexing techniques extend SN. *F-DySNI* [135, 137] converts the sorted list of blocking keys into a braided AVL tree [141] that is updated whenever a query entity arrives. The window is fixed or adaptive, considering as neighbors the nodes exceeding a similarity threshold. F-DySNI is extended in [136] with automatic blocking keys: the weak training set of [72] is coupled with a scoring function that considers both key coverage and block size distribution.

Another group of methods relies on LSH. MinHash LSH is combined with SN in [88]: when searching for the nearest neighbors of a query entity, the entities in large LSH blocks are sorted via a custom scoring function, and then a window of fixed size slides over the sorted list of entities. CF-RDS [68] leverages Hamming LSH, ranking the most similar entities to each query without performing any profile comparison. Instead, it merely aggregates the number of occurrences of each candidate match in the buckets associated with the query entity.

In another line of research, *BlockSketch* [69] organizes the entities inside every block into subblocks according to their similarity. A representative is assigned to each subblock based on its distance from the corresponding blocking key. In this way, every query suffices to be compared with a constant number of entities in the target block in order to detect its most similar entities. *SBlockSketch* [69] adapts this approach to a stream of query entities through an eviction strategy that bounds the number of blocks that need to be maintained in memory.

All these methods are crafted for structured data, assuming a fixed schema of known quality. New techniques are required, though, for the noisy, heterogeneous entities of semistructured data.

Parameter Configuration. Except TB, all Blocking methods involve at least one internal parameter that affects their performance to a large extent [25, 128]. This also affects their relative performance, rendering the selection of the best-performing method into a nontrivial task.

To mitigate this issue, parameter fine-tuning is modeled as an optimization problem in [96]. The configuration space is searched through a genetic algorithm, whose fitness function exploits the labels of some candidate matches. After several generations, the configuration maximizing the fitness function is set as optimal. Yet, this approach involves a large number of parameters itself.

MatchCatcher [86] implements a human-in-the-loop approach combining expert knowledge with labeled instances in order to learn composite blocking schemes. Using string similarity joins, missed duplicates, which share no block, are efficiently detected. To capture them, the expert user adapts the transformation and assignment functions iteratively.

Finally, a method's performance over several labeled datasets is used for fine-tuning its parameters over an unlabeled dataset in [110]. The best choice corresponds to the method that achieves the best combination of runtime and F-Measure across most datasets. However, this is a rather time-consuming approach, given the large number of computations it requires.

None of the above methods satisfies the requirement for automatic, data-driven, a priori parameter configuration of Blocking methods, which thus remains an open problem.

Filtering for Entity Resolution. An interesting direction is to investigate to what extent similarity joins suffice for ER, i.e., representing entity profiles by strings or sets and defining a matching function based on a similarity threshold. Probably, techniques supporting relaxed matching criteria and/or lower-similarity thresholds will be required to achieve high recall, but relatively few Filtering techniques are designed for these cases (see Section 5.3). We also believe that scalability remains an open challenge for string and set similarity joins [50] and that more opportunities exist for transferring ideas and approaches between Blocking and Filtering. Finally, there is a need for extensible, open-source ER tools that incorporate the majority of established Blocking and Filtering methods and apply seamlessly to structured, semistructured, and unstructured data [56].

10 CONCLUSIONS

Efficiency techniques have been an integral part of Entity Resolution since its infancy. We organize the relevant works into Blocking, Filtering, and hybrid techniques, facilitating their understanding and use. We also provide in-depth coverage of each category, further classifying its works into novel subcategories. Lately, the rise of big semistructured data poses challenges to the scalability of efficiency techniques and to their core assumptions: the requirement of Blocking for schema knowledge and of Filtering for high-similarity thresholds. The former led to the introduction of schema-agnostic Blocking and Block Processing techniques, while the latter led to more relaxed criteria of similarity. We cover these new fields in detail, putting in context all relevant works.

REFERENCES

[1] Noha Adly. 2009. Efficient record linkage using a double embedding scheme. In *Proceedings of the 2009 International Conference on Data Mining (DMIN'09)*. 274–281.

[2] Foto Afrati, Anish Das Sarma, David Menestrina, Aditya Parameswaran, and Jeffrey Ullman. 2012. Fuzzy joins using Mapreduce. In *ICDE*. 498–509.

[3] Akiko N. Aizawa and Keizo Oyama. 2005. A fast linkage detection scheme for multi-source information integration. In *Proceedings of the 2005 International Workshop on Challenges in Web Information Retrieval and Integration (WIRI'05)*. 30–39.

[4] Amin Allam, Spiros Skiadopoulos, and Panos Kalnis. 2018. Improved suffix blocking for record linkage and entity resolution. *DKE* 117 (2018), 98–113.

[5] Yasser Altowim, Dmitri V. Kalashnikov, and Sharad Mehrotra. 2014. Progressive approach to relational entity resolution. *PVLDB* 7, 11 (2014), 999–1010.

[6] Yasser Altowim and Sharad Mehrotra. 2017. Parallel progressive approach to entity resolution using Mapreduce. In *Proceedings of the 33rd IEEE International Conference on Data Engineering (ICDE'17)*. 909–920.

[7] Arvind Arasu, Venkatesh Ganti, and Raghav Kaushik. 2006. Efficient exact set-similarity joins. In *Proceedings of the 32nd International Conference on Very Large Data Bases (VLDB'06)*. 918–929.

[8] Samur Araújo, Duc Thanh Tran, Arjen P. de Vries, and Daniel Schwabe. 2015. SERIMI: Class-based matching for instance matching across heterogeneous datasets. *IEEE TKDE* 27, 5 (2015), 1397–1410.

[9] Tiago Brasileiro Araújo, Carlos Eduardo Santos Pires, and Thiago Pereira da Nóbrega. 2017. Spark-based streamlined metablocking. In *Proceedings of the 2017 IEEE Symposium on Computers and Communications (ISCC'17)*. 844–850.

[10] Nikolaus Augsten and Michael H. Böhlen. 2013. *Similarity Joins in Relational Database Systems*. Morgan & Claypool Publishers.

[11] Rohan Baxter, Peter Christen, and Tim Churches. 2003. A comparison of fast blocking methods for record linkage. In *Proceedings of the Workshop on Data Cleaning, Record Linkage and Object Consolidation at the 9th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.

[12] Roberto J. Bayardo, Yiming Ma, and Ramakrishnan Srikant. 2007. Scaling up all pairs similarity search. In *Proceedings of the 16th International Conference on World Wide Web (WWW'07)*. 131–140.

[13] Alexandros Belesiotis, Dimitrios Skoutas, Christodoulos Efstathiades, Vassilis Kaffes, and Dieter Pfoser. 2018. Spatio-textual user matching and clustering based on set similarity joins. *VLDB J.* 27, 3 (2018), 297–320.

[14] Omar Benjelloun, Hector Garcia-Molina, David Menestrina, Qi Su, Steven Euijong Whang, and Jennifer Widom. 2009. Swoosh: A generic approach to entity resolution. *VLDB J.* 18, 1 (2009), 255–276.

[15] Guilherme Dal Bianco, Marcos André Gonçalves, and Denio Duarte. 2018. BLOSS: Effective meta-blocking with almost no effort. *Inf. Syst.* 75 (2018), 75–89.

[16] Mikhail Bilenko, Beena Kamath, and Raymond J. Mooney. 2006. Adaptive blocking: Learning to scale up record linkage. In *Proceedings of the 6th IEEE International Conference on Data Mining (ICDM'06)*. 87–96.

[17] Mikhail Bilenko and Raymond J. Mooney. 2003. Adaptive duplicate detection using learnable string similarity measures. In *Proceedings of the 9th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'03)*. 39–48.

[18] T. Bocek, E. Hunt, and B. Stiller. 2007. *Fast Similarity Search in Large Dictionaries*. Technical Report ifi-2007.02. Department of Informatics, University of Zurich. http://fastss.csg.uzh.ch/.

[19] Panagiotis Bouros, Shen Ge, and Nikos Mamoulis. 2012. Spatio-textual similarity joins. *PVLDB* 6, 1 (2012), 1–12.

[20] Andrei Z. Broder. 1997. On the resemblance and containment of documents. In *Proceedings of the Compression and Complexity of SEQUENCES (SEQUENCES'97)*. 21–29.

[21] Yunbo Cao, Zhiyuan Chen, Jiamin Zhu, Pei Yue, Chin-Yew Lin, and Yong Yu. 2011. Leveraging unlabeled data to scale blocking for record linkage. In *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI'11)*. 2211–2217.

[22] Surajit Chaudhuri, Venkatesh Ganti, and Raghav Kaushik. 2006. A primitive operator for similarity joins in data cleaning. In *Proceedings of the 22nd IEEE International Conference on Data Engineering (ICDE'06)*. 5.

[23] Peter Christen. 2008. Febrl-: An open source data cleaning, deduplication and record linkage system with a graphical user interface. In *Proceedings of the 14th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'08)*. 1065–1068.

[24] Peter Christen. 2012. *Data Matching - Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer.

[25] Peter Christen. 2012. A survey of indexing techniques for scalable record linkage and deduplication. *IEEE TKDE* 24, 9 (2012), 1537–1555.

[26] Peter Christen, Ross W. Gayler, and David Hawking. 2009. Similarity-aware indexing for real-time entity resolution. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management (CIKM'09)*. 1565–1568.

[27] Tobias Christiani and Rasmus Pagh. 2017. Set similarity search beyond MinHash. In *Proceedings of the 49th Annual ACM SIGACT Symposium on Theory of Computing (STOC'17)*. 1094–1107.

[28] Tobias Christiani, Rasmus Pagh, and Johan Sivertsen. 2018. Scalable and robust set similarity join. In *Proceedings of the 34th IEEE International Conference on Data Engineering (ICDE'18)*. 1240–1243.

[29] Vassilis Christophides, Vasilis Efthymiou, and Kostas Stefanidis. 2015. *Entity Resolution in the Web of Data*. Morgan & Claypool Publishers.

[30] Xu Chu, Ihab F. Ilyas, and Paraschos Koutris. 2016. Distributed data deduplication. *PVLDB* 9, 11 (2016), 864–875.

[31] Aaron Clauset, Mark E. J. Newman, and Cristopher Moore. 2004. Finding community structure in very large networks. *Phys. Rev. E* 70, 6 (2004), 066111.

[32] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni. 2004. Locality-sensitive hashing scheme based on p-stable distributions. In *Proceedings of the 20th ACM Symposium on Computational Geometry (SOCG'04)*. 253–262.

[33] Timothy de Vries, Hui Ke, Sanjay Chawla, and Peter Christen. 2009. Robust record linkage blocking using suffix arrays. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management (CIKM'09)*. 305–314.

[34] Timothy de Vries, Hui Ke, Sanjay Chawla, and Peter Christen. 2011. Robust record linkage blocking using suffix arrays and Bloom filters. *TKDD* 5, 2 (2011), 9:1–9:27.

[35] Jeffrey Dean and Sanjay Ghemawat. 2008. MapReduce: Simplified data processing on large clusters. *Commun. ACM* 51, 1 (2008), 107–113.

[36] Dong Deng, Albert Kim, Samuel Madden, and Michael Stonebraker. 2017. SilkMoth: An efficient method for finding related sets with maximum matching constraints. *PVLDB* 10, 10 (2017), 1082–1093.

[37] Dong Deng, Guoliang Li, He Wen, and Jianhua Feng. 2015. An efficient partition based method for exact set similarity joins. *PVLDB* 9, 4 (2015), 360–371.

[38] Xin Luna Dong and Divesh Srivastava. 2015. *Big Data Integration.* Morgan & Claypool Publishers.

[39] Uwe Draisbach and Felix Naumann. 2010. DuDe: The duplicate detection toolkit. In *Proceedings of the 8th International Workshop on Quality in Databases (QDB'10)*.

[40] Uwe Draisbach and Felix Naumann. 2011. A generalization of blocking and windowing algorithms for duplicate detection. In *Proceedings of the 2011 International Conference on Data and Knowledge Engineering (ICDKE'11)*. 18–24.

[41] Uwe Draisbach, Felix Naumann, Sascha Szott, and Oliver Wonneberg. 2012. Adaptive windows for duplicate detection. In *Proceedings of the 28th IEEE International Conference on Data Engineering (ICDE'12)*. 1073–1083.

[42] Songyun Duan, Achille Fokoue, Oktie Hassanzadeh, Anastasios Kementsietsidis, Kavitha Srinivas, and Michael J. Ward. 2012. Instance-based matching of large ontologies using locality-sensitive hashing. In *Proceedings of the 11th International Semantic Web Conference (ISWC'12)*. 49–64.

[43] Muhammad Ebraheem, Saravanan Thirumuruganathan, Shafiq R. Joty, Mourad Ouzzani, and Nan Tang. 2018. Distributed representations of tuples for entity resolution. *PVLDB* 11, 11 (2018), 1454–1467.

[44] Vasilis Efthymiou, George Papadakis, George Papastefanatos, Kostas Stefanidis, and Themis Palpanas. 2015. Parallel meta-blocking: Realizing scalable entity resolution over large, heterogeneous data. In *IEEE Big Data*. 411–420.

[45] Vasilis Efthymiou, George Papadakis, George Papastefanatos, Kostas Stefanidis, and Themis Palpanas. 2017. Parallel meta-blocking for scaling entity resolution over big heterogeneous data. *Inf. Syst.* 65 (2017), 137–157.

[46] Vasilis Efthymiou, Kostas Stefanidis, and Vassilis Christophides. 2015. Big data entity resolution: From highly to somehow similar entity descriptions in the web. In *IEEE Big Data*. 401–410.

[47] Ahmed K. Elmagarmid, Panagiotis G. Ipeirotis, and Vassilios S. Verykios. 2007. Duplicate record detection: A survey. *IEEE TKDE* 19, 1 (2007), 1–16.

[48] Luiz Evangelista, Eli Cortez, Altigran da Silva, and Wagner Meira Jr. 2010. Adaptive and flexible blocking for record linkage tasks. *JIDM* 1, 2 (2010), 167–182.

[49] Ivan P. Fellegi and Alan B. Sunter. 1969. A theory for record linkage. *J. Amer. Statist. Assoc.* 64, 328 (1969), 1183–1210.

[50] Fabian Fier, Nikolaus Augsten, Panagiotis Bouros, Ulf Leser, and Johann-Christoph Freytag. 2018. Set similarity joins on Mapreduce: An experimental survey. *PVLDB* 11, 10 (2018), 1110–1122.

[51] Jeffrey Fisher, Peter Christen, Qing Wang, and Erhard Rahm. 2015. A clustering-based framework to control block sizes for entity resolution. In *Proceedings of the 21st ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'15)*. 279–288.

[52] Dengfeng Gao. 2009. Temporal joins. In *Encyclopedia of Database Systems*. 2982–2987.

[53] Lise Getoor and Ashwin Machanavajjhala. 2012. Entity resolution: Theory, practice & open challenges. *PVLDB* 5, 12 (2012), 2018–2019.

[54] Phan Giang. 2015. A machine learning approach to create blocking criteria for record linkage. *Health Care Manag. Sci.* 18, 1 (2015), 93–105.

[55] Aristides Gionis, Piotr Indyk, and Rajeev Motwani. 1999. Similarity search in high dimensions via hashing. In *Proceedings of 25th International Conference on Very Large Data Bases (VLDB'99)*. 518–529.

[56] Behzad Golshan, Alon Y. Halevy, George A. Mihaila, and Wang-Chiew Tan. 2017. Data integration: After the teenage years. In *ACM PODS*. 101–106.

[57] Luis Gravano, Panagiotis G. Ipeirotis, H. V. Jagadish, Nick Koudas, S. Muthukrishnan, and Divesh Srivastava. 2001. Approximate string joins in a database (Almost) for free. In *Proceedings of 27th International Conference on Very Large Data Bases (VLDB'01)*. 491–500.

[58] Thomas Gschwind, Christoph Miksovic, Katsiaryna Mirylenka, and Paolo Scotton. 2019. Fast record linkage for company entities. *CoRR* abs/1907.08667 (2019).

[59] Lifang Gu and Rohan A. Baxter. 2004. Adaptive filtering for efficient record linkage. In *Proceedings of the 4th SIAM International Conference on Data Mining (SDM'04)*. 477–481.

[60] Mauricio A. Hernández and Salvatore J. Stolfo. 1995. The merge/purge problem for large databases. In *Proceedings of the 1995 ACM SIGMOD International Conference on Management of Data (SIGMOD'95)*. 127–138.

[61] Mauricio A. Hernández and Salvatore J. Stolfo. 1998. Real-world data is dirty: Data cleansing and the merge/purge problem. *Data Min. Knowl. Discov.* 2, 1 (1998), 9–37.

[62] Robert Isele, Anja Jentzsch, and Christian Bizer. 2011. Efficient multidimensional blocking for link discovery without losing recall. In *Proceedings of the 14th International Workshop on the Web and Database (WebDB'11)*.

[63] Edwin H. Jacox and Hanan Samet. 2007. Spatial join techniques. *ACM TODS* 32, 1 (2007), 7.

[64] Yu Jiang, Guoliang Li, Jianhua Feng, and Wen-Syan Li. 2014. String similarity joins: An experimental evaluation. *PVLDB* 7, 8 (2014), 625–636.

[65] Liang Jin, Chen Li, and Sharad Mehrotra. 2003. Efficient record linkage in large data sets. In *Proceedings of the 8th International Conference on Database Systems for Advanced Applications (DASFAA'03)*. 137–146.

[66] Pawel Jurczyk, James J. Lu, Li Xiong, Janet D. Cragan, and Adolfo Correa. 2008. Fine-grained record integration and linkage tool. *Birth Defects Res. A: Clin. Mol. Teratol.* 82, 11 (2008), 822–829.

[67] Murat Kantarcioglu, Ali Inan, Wei Jiang, and Bradley Malin. 2009. Formal anonymity models for efficient privacy-preserving joins. *DKE* 68, 11 (2009), 1206–1223.

[68] Dimitrios Karapiperis, Aris Gkoulalas-Divanis, and Vassilios S. Verykios. 2018. Fast schemes for online record linkage. *Data Min. Knowl. Discov.* 32, 5 (2018), 1229–1250.

[69] Dimitrios Karapiperis, Aris Gkoulalas-Divanis, and Vassilios S. Verykios. 2018. Summarization algorithms for record linkage. In *Proceedings of the 21st International Conference on Extending Database Technology (EDBT'18)*. 73–84.

[70] Dimitrios Karapiperis, Dinusha Vatsalan, Vassilios S. Verykios, and Peter Christen. 2016. Efficient record linkage using a compact hamming space. In *Proceedings of the 19th International Conference on Extending Database Technology (EDBT'16)*. 209–220.

[71] Dimitrios Karapiperis and Vassilios S. Verykios. 2016. A fast and efficient hamming LSH-based scheme for accurate linkage. *KAIS* 49, 3 (2016), 861–884.

[72] Mayank Kejriwal and Daniel P. Miranker. 2013. An unsupervised algorithm for learning blocking schemes. In *Proceedings of the 13th IEEE International Conference on Data Mining (ICDM'13)*. 340–349.

[73] Mayank Kejriwal and Daniel P. Miranker. 2014. A two-step blocking scheme learner for scalable link discovery. In *Proceedings of the 9th International Workshop on Ontology Matching (OM'14)*. 49–60.

[74] Mayank Kejriwal and Daniel P. Miranker. 2015. A DNF blocking scheme learner for heterogeneous datasets. *CoRR* abs/1501.01694 (2015).

[75] Batya Kenig and Avigdor Gal. 2013. MFIBlocks: An effective blocking algorithm for entity resolution. *Inf. Syst.* 38, 6 (2013), 908–926.

[76] Hung-sik Kim and Dongwon Lee. 2010. HARRA: Fast iterative hashed record linkage for large-scale data collections. In *Proceedings of the 13th International Conference on Extending Database Technology (EDBT'10)*. 525–536.

[77] Lars Kolb, Andreas Thor, and Erhard Rahm. 2012. Dedoop: Efficient deduplication with Hadoop. *PVLDB* 5, 12 (2012), 1878–1881.

[78] Lars Kolb, Andreas Thor, and Erhard Rahm. 2012. Load balancing for Mapreduce-based entity resolution. In *Proceedings of the 28th IEEE International Conference on Data Engineering (ICDE'12)*. 618–629.

[79] Lars Kolb, Andreas Thor, and Erhard Rahm. 2012. Multi-pass sorted neighborhood blocking with MapReduce. *Comput. Sci. R&D* 27, 1 (2012), 45–63.

[80] Pradap Konda, Sanjib Das, Paul Suganthan G. C., AnHai Doan, Adel Ardalan, Jeffrey R. Ballard, Han Li, Fatemah Panahi, Haojun Zhang, Jeffrey F. Naughton, Shishir Prasad, Ganesh Krishnan, Rohit Deep, and Vijay Raghavendra. 2016. Magellan: Toward building entity matching management systems. *PVLDB* 9, 12 (2016), 1197–1208.

[81] Hanna Köpcke, Andreas Thor, and Erhard Rahm. 2009. Comparative evaluation of entity resolution approaches with FEVER. *PVLDB* 2, 2 (2009), 1574–1577.

[82] Hanna Köpcke, Andreas Thor, and Erhard Rahm. 2010. Evaluation of entity resolution approaches on real-world match problems. *PVLDB* 3, 1 (2010), 484–493.

[83] Chen Li, Jiaheng Lu, and Yiming Lu. 2008. Efficient merging and filtering algorithms for approximate string searches. In *Proceedings of the 24th IEEE International Conference on Data Engineering (ICDE'08)*. 257–266.

[84] Guoliang Li, Dong Deng, Jiannan Wang, and Jianhua Feng. 2011. PASS-JOIN: A partition-based method for similarity joins. *PVLDB* 5, 3 (2011), 253–264.

[85] Guoliang Li, Jian He, Dong Deng, and Jian Li. 2015. Efficient similarity join and search on multi-attribute data. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data (SIGMOD'15)*. 1137–1151.

[86] Han Li, Pradap Konda, Paul Suganthan, AnHai Doan, Benjamin Snyder, Youngchoon Park, Ganesh Krishnan, Rohit Deep, and Vijay Raghavendra. 2018. MatchCatcher: A debugger for blocking in entity matching. In *Proceedings of the 21st International Conference on Extending Database Technology (EDBT'18)*. 193–204.

[87] Yaping Li and Minghua Chen. 2008. Privacy preserving joins. In *Proceedings of the 24th IEEE International Conference on Data Engineering (ICDE'08)*. 1352–1354.

[88] Huizhi Liang, Yanzhe Wang, Peter Christen, and Ross W. Gayler. 2014. Noise-tolerant approximate blocking for dynamic real-time entity resolution. In *Proceedings of the 18th Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD'14)*. 449–460.

[89] Wei Lu, Xiaoyong Du, Marios Hadjieleftheriou, and Beng Chin Ooi. 2014. Efficiently supporting edit distance based string similarity search using B⁺-trees. *IEEE TKDE* 26, 12 (2014), 2983–2996.

[90] Qin Lv, William Josephson, Zhe Wang, Moses Charikar, and Kai Li. 2007. Multi-probe LSH: Efficient indexing for high-dimensional similarity search. In *Proceedings of the 33rd International Conference on Very Large Data Bases (VLDB'07)*. 950–961.

[91] Kun Ma, Fusen Dong, and Bo Yang. 2015. Large-scale schema-free data deduplication approach with adaptive sliding window using mapreduce. *Comput. J.* 58, 11 (2015), 3187–3201.

[92] Yongtao Ma and Thanh Tran. 2013. TYPiMatch: Type-specific unsupervised learning of keys and key values for heterogeneous web data integration. In *Proceedings of the 6th ACM International Conference on Web Search and Data Mining (WSDM'13)*. 325–334.

[93] Pankaj Malhotra, Puneet Agarwal, and Gautam Shroff. 2014. Graph-parallel entity resolution using LSH & IMM. In *Proceedings of the Workshops of the EDBT/ICDT 2014 Joint Conference*. 41–49.

[94] Willi Mann and Nikolaus Augsten. 2014. PEL: Position-enhanced length filter for set similarity joins. In *Grundlagen Datenbanken*. 89–94.

[95] Willi Mann, Nikolaus Augsten, and Panagiotis Bouros. 2016. An empirical evaluation of set similarity join techniques. *PVLDB* 9, 9 (2016), 636–647.

[96] Ruhaila Maskat, Norman W. Paton, and Suzanne M. Embury. 2016. Pay-as-you-go configuration of entity resolution. *T-LSD-KCS* 29 (2016), 40–65.

[97] Andrew McCallum, Kamal Nigam, and Lyle H. Ungar. 2000. Efficient clustering of high-dimensional data sets with application to reference matching. In *Proceedings of the Workshops of the EDBT/ICDT 2014 Joint Conference*. 169–178.

[98] W. P. McNeill, Hakan Kardes, and Andrew Borthwick. 2012. Dynamic record blocking: Efficient linking of massive databases in Mapreduce. In *Proceedings of the 10th International Workshop on Quality in Databases (QDB'12)*.

[99] Demetrio Gomes Mestre, Carlos Eduardo S. Pires, and Dimas C. Nascimento. 2015. Adaptive sorted neighborhood blocking for entity matching with Mapreduce. In *Proceedings of the 30th Annual ACM Symposium on Applied Computing (SAC'15)*. 981–987.

[100] Matthew Michelson and Craig A. Knoblock. 2006. Learning blocking schemes for record linkage. In *Proceedings of the 21st National Conference on Artificial Intelligence (AAAI'06)*. 440–445.

[101] Tomas Mikolov, Ilya Sutskever, Kai Chen, Gregory S. Corrado, and Jeffrey Dean. 2013. Distributed representations of words and phrases and their compositionality. In *Proceedings of the 27th Annual Conference on Neural Information Processing Systems (NIPS'13)*. 3111–3119.

[102] Dimas Cassimiro Nascimento, Carlos Eduardo Santos Pires, and Demetrio Gomes Mestre. 2020. Exploiting block co-occurrence to control block sizes for entity resolution. *Knowl. Inf. Syst.* 62, 1 (2020), 359–400.

[103] Sahand Negahban, Benjamin I. P. Rubinstein, and Jim Gemmell. 2012. Scaling multiple-source entity resolution using statistically efficient transfer learning. In *Proceedings of the 21st ACM International Conference on Information and Knowledge Management (CIKM'12)*. 2224–2228.

[104] E. D. Nelson and J. R. Talburt. 2011. Entity resolution for longitudinal studies in education using OYSTER. In *Proceedings of the 2011 International Conference on Information and Knowledge Engineering (IKE'11)*.

[105] Markus Nentwig, Michael Hartung, Axel Ngomo, and Erhard Rahm. 2017. A survey of current link discovery frameworks. *Semantic Web* 8, 3 (2017), 419–436.

[106] Axel Ngomo. 2013. ORCHID - Reduction-ratio-optimal computation of geo-spatial distances for link discovery. In *Proceedings of the 12th International Semantic Web Conference (ISWC'12)*. 395–410.

[107] Axel Ngomo and Sören Auer. 2011. LIMES - A time-efficient approach for large-scale link discovery on the web of data. In *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI'11)*. 2312–2317.

[108] Andriy Nikolov, Victoria Uren, and Enrico Motta. 2007. KnoFuss: A comprehensive architecture for knowledge fusion. In *Proceedings of the 4th International Conference on Knowledge Capture (K-CAP'07)*. 185–186.

[109] Jordi Nin, Victor Muntés-Mulero, Norbert Martínez-Bazan, and Josep-Lluís Larriba-Pey. 2007. On the use of semantic blocking techniques for data cleansing and integration. In *Proceedings of the 11th International Database Engineering and Applications Symposium (IDEAS'07)*. 190–198.

[110] Kevin O'Hare, Anna Jurek, and Cassio de Campos. 2018. A new technique of selecting an optimal blocking method for better record linkage. *Inf. Syst.* 77 (2018), 151–166.

[111] Kevin O'Hare, Anna Jurek-Loughrey, and Cassio de Campos. 2019. A review of unsupervised and semi-supervised blocking methods for record linkage. In *Linking and Mining Heterogeneous and Multi-view Data*. 79–105.

[112] George Papadakis, George Alexiou, George Papastefanatos, and Georgia Koutrika. 2015. Schema-agnostic vs schema-based configurations for blocking methods on homogeneous data. *PVLDB* 9, 4 (2015), 312–323.

[113] George Papadakis, Konstantina Bereta, Themis Palpanas, and Manolis Koubarakis. 2017. Multi-core meta-blocking for big linked data. In *Proceedings of the 13th International Conference on Semantic Systems (SEMANTICS'17)*. 33–40.

[114] George Papadakis, Gianluca Demartini, Peter Fankhauser, and Philipp Kärger. 2010. The missing links: Discovering hidden same-as links among a billion of triples. In *Proceedings of the 12th International Conference on Information Integration and Web-based Applications and Services (iiWAS'10)*. 453–460.

[115] George Papadakis, George Giannakopoulos, Claudia Niederée, Themis Palpanas, and Wolfgang Nejdl. 2011. Detecting and exploiting stability in evolving heterogeneous information spaces. In *Proceedings of the 2011 Joint International Conference on Digital Libraries (JCDL'11)*. 95–104.

[116] George Papadakis, Ekaterini Ioannou, Claudia Niederée, and Peter Fankhauser. 2011. Efficient entity resolution for large heterogeneous information spaces. In *Proceedings of the 4th International Conference on Web Search and Web Data Mining (WSDM'11)*. 535–544.

[117] George Papadakis, Ekaterini Ioannou, Claudia Niederée, Themis Palpanas, and Wolfgang Nejdl. 2011. Eliminating the redundancy in blocking-based entity resolution methods. In *Proceedings of the 2011 Joint International Conference on Digital Libraries (JCDL'11)*. 85–94.

[118] George Papadakis, Ekaterini Ioannou, Claudia Niederée, Themis Palpanas, and Wolfgang Nejdl. 2011. To compare or not to compare: Making entity resolution more efficient. In *Proceedings of the International Workshop on Semantic Web Information Management (SWIM'11)*. 3.

[119] George Papadakis, Ekaterini Ioannou, Claudia Niederée, Themis Palpanas, and Wolfgang Nejdl. 2012. Beyond 100 million entities: Large-scale blocking-based resolution for heterogeneous data. In *Proceedings of the 5th International Conference on Web Search and Web Data Mining (WSDM'12)*. 53–62.

[120] George Papadakis, Ekaterini Ioannou, Themis Palpanas, Claudia Niederée, and Wolfgang Nejdl. 2013. A blocking framework for entity resolution in highly heterogeneous information spaces. *IEEE TKDE* 25, 12 (2013), 2665–2682.

[121] George Papadakis, Georgia Koutrika, Themis Palpanas, and Wolfgang Nejdl. 2014. Meta-blocking: Taking entity resolution to the next level. *IEEE TKDE* 26, 8 (2014), 1946–1960.

[122] George Papadakis and Wolfgang Nejdl. 2011. Efficient entity resolution methods for heterogeneous information spaces. In *Workshops Proceedings of the 27th IEEE International Conference on Data Engineering*. 304–307.

[123] George Papadakis and Themis Palpanas. 2016. Blocking for large-scale entity resolution: Challenges, algorithms, and practical examples. In *Proceedings of the 32nd IEEE International Conference on Data Engineering (ICDE'16)*. 1436–1439.

[124] George Papadakis and Themis Palpanas. 2018. Web-scale, schema-agnostic, end-to-end entity resolution. In *Companion Volume of The Web Conference 2018 (WWW'18)*.

[125] George Papadakis, George Papastefanatos, and Georgia Koutrika. 2014. Supervised meta-blocking. *PVLDB* 7, 14 (2014), 1929–1940.

[126] George Papadakis, George Papastefanatos, Themis Palpanas, and Manolis Koubarakis. 2016. Boosting the efficiency of large-scale entity resolution with enhanced meta-blocking. *Big Data Res.* 6 (2016), 43–63.

[127] George Papadakis, George Papastefanatos, Themis Palpanas, and Manolis Koubarakis. 2016. Scaling entity resolution to large, heterogeneous data with enhanced meta-blocking. In *Proceedings of the 19th International Conference on Extending Database Technology (EDBT'16)*. 221–232.

[128] George Papadakis, Jonathan Svirsky, Avigdor Gal, and Themis Palpanas. 2016. Comparative analysis of approximate blocking techniques for entity resolution. *PVLDB* 9, 9 (2016), 684–695.

[129] George Papadakis, Leonidas Tsekouras, Emmanouil Thanos, George Giannakopoulos, Themis Palpanas, and Manolis Koubarakis. 2018. The return of JedAI: End-to-end entity resolution for structured and semi-structured data. *PVLDB* 11, 12 (2018), 1950–1953.

[130] Thorsten Papenbrock, Arvid Heise, and Felix Naumann. 2015. Progressive duplicate detection. *IEEE TKDE* 27, 5 (2015), 1316–1329.

[131] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. Glove: Global vectors for word representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP'14)*. 1532–1543.

[132] Sven Puhlmann, Melanie Weis, and Felix Naumann. 2006. XML Duplicate detection using sorted neighborhoods. In *Proceedings of the 10th International Conference on Extending Database Technology (EDBT'06)*. 773–791.

[133] Jianbin Qin, Wei Wang, Yifei Lu, Chuan Xiao, and Xuemin Lin. 2011. Efficient exact edit similarity query processing with the asymmetric signature scheme. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data (SIGMOD'11)*. 1033–1044.

[134] Jianbin Qin and Chuan Xiao. 2018. Pigeonring: A principle for faster thresholded similarity search. *PVLDB* 12, 1 (2018), 28–42.

[135] Banda Ramadan and Peter Christen. 2014. Forest-based dynamic sorted neighborhood indexing for real-time entity resolution. In *Proceedings of the 23rd ACM International Conference on Information and Knowledge Management (CIKM'14)*. 1787–1790.

[136] Banda Ramadan and Peter Christen. 2015. Unsupervised blocking key selection for real-time entity resolution. In *Proceedings of the 19th Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD'15)*. 574–585.

[137] Banda Ramadan, Peter Christen, Huizhi Liang, and Ross W. Gayler. 2015. Dynamic sorted neighborhood indexing for real-time entity resolution. *J. Data Inf. Qual.* 6, 4, Article 15 (2015), 15:1–15:29 pages.

[138] Banda Ramadan, Peter Christen, Huizhi Liang, Ross W. Gayler, and David Hawking. 2013. Dynamic similarity-aware inverted indexing for real-time entity resolution. In *International Workshops of the 17th Pacific-Asia Conference on Knowledge Discovery and Data Mining*. 47–58.

[139] Thilina Ranbaduge, Dinusha Vatsalan, and Peter Christen. 2016. Scalable block scheduling for efficient multi-database record linkage. In *Proceedings of the 16th IEEE International Conference on Data Mining (ICDM'16)*. 1161–1166.

[140] Leonardo Andrade Ribeiro and Theo Härder. 2011. Generalizing prefix filtering to improve set similarity joins. *Inf. Syst.* 36, 1 (2011), 62–78.

[141] Stephen V. Rice. 2007. Braided AVL trees for efficient event sets and ranked sets in the SIMSCRIPT III simulation programming language. In *Western MultiConference on Computer Simulation*. 150–155.

[142] Chuitian Rong, Chunbin Lin, Yasin N. Silva, Jianguo Wang, Wei Lu, and Xiaoyong Du. 2017. Fast and scalable distributed set similarity joins for big data analytics. In *Proceedings of the 33rd IEEE International Conference on Data Engineering (ICDE'17)*. 1059–1070.

[143] Chuitian Rong, Wei Lu, Xiaoli Wang, Xiaoyong Du, Yueguo Chen, and Anthony K. H. Tung. 2013. Efficient and scalable processing of string similarity join. *IEEE TKDE* 25, 10 (2013), 2217–2230.

[144] Sunita Sarawagi and Alok Kirpal. 2004. Efficient set joins on similarity predicates. In *Proceedings of the 2004 ACM SIGMOD International Conference on Management of Data (SIGMOD'04)*. 743–754.

[145] Murat Sariyar, Andreas Borg, and Klaus Pommerening. 2011. Controlling false match rates in record linkage using extreme value theory. *J. Biomed. Inf.* 44, 4 (2011), 648–654.

[146] Anish Das Sarma, Ankur Jain, Ashwin Machanavajjhala, and Philip Bohannon. 2012. An automatic blocking mechanism for large-scale de-duplication tasks. In *Proceedings of the 21st ACM International Conference on Information and Knowledge Management (CIKM'12)*. 1055–1064.

[147] Venu Satuluri and Srinivasan Parthasarathy. 2012. Bayesian locality sensitive hashing for fast similarity search. *PVLDB* 5, 5 (2012), 430–441.

[148] Wei Shen, Jianyong Wang, and Jiawei Han. 2015. Entity linking with a knowledge base: Issues, techniques, and solutions. *IEEE TKDE* 27, 2 (2015), 443–460.

[149] Liangcai Shu, Aiyou Chen, Ming Xiong, and Weiyi Meng. 2011. Efficient spectral neighborhood blocking for entity resolution. In *Proceedings of the 27th International Conference on Data Engineering (ICDE'11)*. 1067–1078.

[150] Giovanni Simonini, Sonia Bergamaschi, and H. V. Jagadish. 2016. BLAST: A loosely schema-aware meta-blocking approach for entity resolution. *PVLDB* 9, 12 (2016), 1173–1184.

[151] Giovanni Simonini, Luca Gagliardelli, Sonia Bergamaschi, and H. V. Jagadish. 2019. Scaling entity resolution: A loosely schema-aware approach. *Inf. Syst.* 83 (2019), 145–165.

[152] Giovanni Simonini, George Papadakis, Themis Palpanas, and Sonia Bergamaschi. 2019. Schema-agnostic progressive entity resolution. *IEEE TKDE* 31, 6 (2019), 1208–1221.

[153] Dezhao Song. 2012. Scalable and domain-independent entity coreference: Establishing high quality data linkages across heterogeneous data sources. In *Proceedings of the 11th International Semantic Web Conference (ISWC'12)*. 424–432.

[154] Dezhao Song and Jeff Heflin. 2011. Automatically generating data linkages using a domain-independent candidate selection approach. In *Proceedings of the 10th International Semantic Web Conference (ISWC'11)*. 649–664.

[155] Dezhao Song, Yi Luo, and Jeff Heflin. 2017. Linking heterogeneous data in the semantic web using scalable and domain-independent candidate selection. *IEEE TKDE* 29, 1 (2017), 143–156.

[156] Kostas Stefanidis, Vassilis Christophides, and Vasilis Efthymiou. 2017. Web-scale blocking, iterative and progressive entity resolution. In *Proceedings of the 33rd IEEE International Conference on Data Engineering (ICDE'17)*. 1459–1462.

[157] Kostas Stefanidis, Vasilis Efthymiou, Melanie Herschel, and Vassilis Christophides. 2014. Entity resolution in the web of data. In *Companion Volume of the 23rd International World Wide Web Conference (WWW'14)*. 203–204.

[158] Rebecca C. Steorts, Samuel L. Ventura, Mauricio Sadinle, and Stephen E. Fienberg. 2014. A comparison of blocking methods for record linkage. In *Privacy in Statistical Databases*. 253–268.

[159] Paul Suganthan, Adel Ardalan, AnHai Doan, and Aditya Akella. 2018. Smurf: Self-service string matching using random forests. *PVLDB* 12, 3 (2018), 278–291.

[160] Ji Sun, Zeyuan Shang, Guoliang Li, Zhifeng Bao, and Dong Deng. 2019. Balance-aware distributed string similarity-based query processing system. *PVLDB* 12, 9 (2019), 961–974.

[161] Ji Sun, Zeyuan Shang, Guoliang Li, Dong Deng, and Zhifeng Bao. 2017. Dima: A distributed in-memory similarity-based query processing system. *PVLDB* 10, 12 (2017), 1925–1928.

[162] Wenbo Tao, Dong Deng, and Michael Stonebraker. 2017. Approximate string joins with abbreviations. *PVLDB* 11, 1 (2017), 53–65.

[163] Yufei Tao, Ke Yi, Cheng Sheng, and Panos Kalnis. 2009. Quality and efficiency in high dimensional nearest neighbor search. In *Proceedings of the ACM SIGMOD International Conference on Management of Data (SIGMOD'09)*. 563–576.

[164] Saravanan Thirumuruganathan, Shameem Ahamed Puthiya Parambath, Mourad Ouzzani, Nan Tang, and Shafiq R. Joty. 2018. Reuse and adaptation for entity resolution through transfer learning. *CoRR* abs/1809.11084 (2018).

[165] Dinusha Vatsalan, Peter Christen, and Vassilios S. Verykios. 2013. A taxonomy of privacy-preserving record linkage techniques. *Inf. Syst.* 38, 6 (2013), 946–969.

[166] Rares Vernica, Michael J. Carey, and Chen Li. 2010. Efficient parallel set-similarity joins using MapReduce. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of Data (SIGMOD'10)*. 495–506.

[167] Julius Volz, Christian Bizer, Martin Gaedke, and Georgi Kobilarov. 2009. Silk-a link discovery framework for the web of data. In *Proceedings of the WWW2009 Workshop on Linked Data on the Web (LDOW'09)*. 53 pages.

[168] Jiannan Wang, Guoliang Li, and Jianhua Feng. 2010. Trie-join: Efficient trie-based string similarity joins with edit-distance constraints. *PVLDB* 3, 1 (2010), 1219–1230.

[169] Jiannan Wang, Guoliang Li, and Jianhua Feng. 2012. Can we beat the prefix filtering?: An adaptive framework for similarity join and search. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data (SIGMOD'12)*. 85–96.

[170] Jiannan Wang, Guoliang Li, and Jianhua Feng. 2014. Extending string similarity join to tolerant fuzzy token matching. *ACM TODS* 39, 1 (2014), 7:1–7:45.

[171] Jin Wang, Chunbin Lin, and Carlo Zaniolo. 2019. MF-Join: Efficient fuzzy string similarity join with multi-level filtering. In *Proceedings of the 35th IEEE International Conference on Data Engineering (ICDE'19)*. 386–397.

[172] Jingdong Wang, Heng Tao Shen, Jingkuan Song, and Jianqiu Ji. 2014. Hashing for similarity search: A survey. *CoRR* abs/1408.2927 (2014).

[173] Jiaying Wang, Xiaochun Yang, Bin Wang, and Chengfei Liu. 2017. LS-Join: Local similarity join on string collections. *IEEE TKDE* 29, 9 (2017), 1928–1942.

[174] Pei Wang, Chuan Xiao, Jianbin Qin, Wei Wang, Xiaoyang Zhang, and Yoshiharu Ishikawa. 2016. Local similarity search for unstructured text. In *Proceedings of the 2016 ACM International Conference on Management of Data (SIGMOD'16)*. 1991–2005.

[175] Qing Wang, Mingyuan Cui, and Huizhi Liang. 2016. Semantic-aware blocking for entity resolution. *IEEE TKDE* 28, 1 (2016), 166–180.

[176] Wei Wang, Jianbin Qin, Chuan Xiao, Xuemin Lin, and Heng Tao Shen. 2013. VChunkJoin: An efficient algorithm for edit similarity joins. *IEEE TKDE* 25, 8 (2013), 1916–1929.

[177] Xubo Wang, Lu Qin, Xuemin Lin, Ying Zhang, and Lijun Chang. 2017. Leveraging set relations in exact set similarity join. *PVLDB* 10, 9 (2017), 925–936.

[178] Steven Euijong Whang, David Marmaros, and Hector Garcia-Molina. 2013. Pay-as-you-go entity resolution. *IEEE TKDE* 25, 5 (2013), 1111–1124.

[179] Steven Euijong Whang, David Menestrina, Georgia Koutrika, Martin Theobald, and Hector Garcia-Molina. 2009. Entity resolution with iterative blocking. In *SIGMOD*. 219–232.

[180] Chuan Xiao, Wei Wang, and Xuemin Lin. 2008. Ed-Join: An efficient algorithm for similarity joins with edit distance constraints. *PVLDB* 1, 1 (2008), 933–944.

[181] Chuan Xiao, Wei Wang, Xuemin Lin, and Haichuan Shang. 2009. Top-k set similarity joins. In *ICDE*. 916–927.

[182] Chuan Xiao, Wei Wang, Xuemin Lin, and Jeffrey Xu Yu. 2008. Efficient similarity joins for near duplicate detection. In *WWW*. 131–140.

[183] Chuan Xiao, Wei Wang, Xuemin Lin, Jeffrey Xu Yu, and Guoren Wang. 2011. Efficient similarity joins for near-duplicate detection. *ACM TODS* 36, 3 (2011), 15:1–15:41.

[184] Pengfei Xu and Jiaheng Lu. 2019. Towards a unified framework for string similarity joins. *PVLDB* 12, 11 (2019), 1289–1302.

[185] Su Yan, Dongwon Lee, Min-Yen Kan, and C. Lee Giles. 2007. Adaptive sorted neighborhood methods for efficient record linkage. In *JCDL*. 185–194.

[186] Wei Yan, Yuan Xue, and Bradley Malin. 2013. Scalable load balancing for Mapreduce-based record linkage. In *IPCCC*. 1–10.

[187] Minghe Yu, Guoliang Li, Dong Deng, and Jianhua Feng. 2016. String similarity search and join: A survey. *Frontiers Comput. Sci.* 10, 3 (2016), 399–417.

[188] Minghe Yu, Jin Wang, Guoliang Li, Yong Zhang, Dong Deng, and Jianhua Feng. 2017. A unified framework for string similarity search with edit-distance constraint. *VLDB J.* 26, 2 (2017), 249–274.

[189] Xingliang Yuan, Xinyu Wang, Cong Wang, Chenyun Yu, and Sarana Nutanong. 2017. Privacy-preserving similarity joins over encrypted data. *IEEE TIFS* 12, 11 (2017), 2763–2775.

[190] Jiaqi Zhai, Yin Lou, and Johannes Gehrke. 2011. ATLAS: A probabilistic algorithm for high dimensional similarity search. In *SIGMOD*. 997–1008.

[191] Fulin Zhang, Zhipeng Gao, and Kun Niu. 2017. A pruning algorithm for meta-blocking based on cumulative weight. *J. Phys. Conf. Ser.* 887, 1 (2017), 012058. https://iopscience.iop.org/article/10.1088/1742-6596/887/1/012058.

[192] Yong Zhang, Xiuxing Li, Jin Wang, Ying Zhang, Chunxiao Xing, and Xiaojie Yuan. 2017. An efficient framework for exact set similarity search using tree structure indexes. In *ICDE*. 759–770.

[193] Yong Zhang, Jiacheng Wu, Jin Wang, and Chunxiao Xing. 2020. A transformation-based framework for knn set similarity search. *IEEE TKDE* 32, 3 (2020), 409–423.

[194] Zhenjie Zhang, Marios Hadjieleftheriou, Beng Chin Ooi, and Divesh Srivastava. 2010. Bed-tree: An all-purpose index structure for string similarity search based on edit distance. In *SIGMOD*. 915–926.

[195] Erkang Zhu, Dong Deng, Fatemeh Nargesian, and Renée J. Miller. 2019. JOSIE: Overlap set similarity search for finding joinable tables in data lakes. In *SIGMOD*. 847–864.